


## Problem 1

a)



Given :-

- Learning Rate ( $\alpha$ ) = 0.1

- Inputs  $x_1 = 0$ ,  $x_2 = 1$

- Output  $y = 1$

- All weights & biases = 1

i.e.  $w_1 = w_2 = w_3 = w_4 = w_5 = w_6 = w_7 = w_8 = w_9 = b_1 =$

$b_2 = b_3 = b_4 = 1$

Step I Forward Pass :- Using ReLU  $\Rightarrow f(x) = \begin{cases} 0 & \text{if } x < 0 \\ x & \text{if } x \geq 0 \end{cases}$  &  $f(x) = \max(0, x)$

- For  $n_1$

$$z_1 = x_1 w_1 + x_2 w_2 + b_1 = 0 \cdot 1 + 1 \cdot 1 + 1 = 2$$

$$a_1 = \text{ReLU}(z_1) = \text{ReLU}(2) = \underline{2} \quad \text{as it is +ve} \\ \text{ \& } f(2) = \max(0, 2) = 2$$

- For  $n_2$

$$z_2 = x_1 w_3 + x_2 w_4 + b_2 = 0 \cdot 1 + 1 \cdot 1 + 1 = 2$$

$$a_2 = \text{ReLU}(z_2) = \text{ReLU}(2) = \underline{2}$$

• For  $n_3$

$$Z_3 = x_1 \cdot w_5 + x_2 \cdot w_6 + b_3 = 0.1 + 1.1 + 1 = 2$$

$$a_3 = \text{ReLU}(Z_3) = \text{ReLU}(2) = 2 \quad [\text{as } \text{ReLU}(x) = \max(0, x)]$$

Used ReLU for hidden units and now using sigmoid for output node

$$\begin{aligned} Z_4 &= w_7 \cdot a_1 + w_8 \cdot a_2 + w_9 \cdot a_3 + b_4 \\ &= 1 \cdot 2 + 1 \cdot 2 + 1 \cdot 2 + 1 \\ &= 7 \end{aligned}$$

Using sigmoid Activation

$$\sigma(x) = \frac{1}{1+e^{-x}}$$

$$\text{output } (\hat{y}) = \sigma(Z_4) = \frac{1}{1+e^{-7}} \quad [\text{Using Sigmoid Activation}]$$

$$= \frac{1}{1+0.000911}$$

$$= \frac{1}{1.000911}$$

$$= 0.9990889488$$

$$\hat{y} \approx 0.9991$$

Step II Error

$$\text{Loss Function} = \text{MSE} = \frac{1}{2} (y - \hat{y})^2$$

$$\text{MSE} = \frac{1}{2} (1 - 0.9991)^2 = \frac{1}{2} (0.0009)^2$$

$$= 4.05 \times 10^{-7}$$



### Step III Backpropagation

$$L = \frac{1}{2} (y - \hat{y})^2$$

derivative w.r.t  $\hat{y}$

$$\frac{\partial L}{\partial \hat{y}} = \frac{\partial}{\partial \hat{y}} \left( \frac{1}{2} (y - \hat{y})^2 \right) = \frac{1}{2} \cdot 2(y - \hat{y}) \cdot \frac{\partial (y - \hat{y})}{\partial \hat{y}}$$

with chain Rule

$$\frac{\partial L}{\partial \hat{y}} = -(y - \hat{y}) = \underline{\underline{\hat{y} - y}}$$

where we have  $y$  (ground truth) = 1 &  $\hat{y} = 0.9991$

$\therefore$  we get,

$$\frac{\partial L}{\partial \hat{y}} = 0.9991 - 1 = \underline{\underline{-0.0009}}$$

• Gradient of Loss w.r.t.  $z_4$

$$\frac{\partial L}{\partial z_4} = \frac{\partial L}{\partial \hat{y}} \cdot \sigma'(z_4)$$

where

$$\sigma'(z_4) = \sigma(z_4) \cdot (1 - \sigma(z_4)) = 0.9991 \cdot (1 - 0.9991) \approx 0.0009$$

$$\therefore \frac{\partial L}{\partial z_4} = -0.0009 \times 0.0009 = -8.1 \times 10^{-7}$$



• Gradient wrt  $w_7, w_8, w_9$  are

- For  $w_7$ :

$$\frac{\partial L}{\partial w_7} = \frac{\partial L}{\partial z_4} \cdot a_1 = -8.1 \times 10^{-7} \times 2 = -1.62 \times 10^{-6}$$

- For  $w_8$ :

$$\frac{\partial L}{\partial w_8} = \frac{\partial L}{\partial z_4} \cdot a_2 = -8.1 \times 10^{-7} \times 2 = -1.62 \times 10^{-6}$$

- For  $w_9$ :

$$\frac{\partial L}{\partial w_9} = \frac{\partial L}{\partial z_4} \cdot a_3 = -8.1 \times 10^{-7} \times 2 = -1.6 \times 10^{-6}$$

To compute updated weights using formula

$$w' = w - \alpha \frac{\partial L}{\partial w}$$

we have  $\alpha = 0.1$

$\therefore$  for  $w_7$ :

$$w_7' = w_7 - \alpha \frac{\partial L}{\partial w_7} = 1 - 0.1(-1.62 \times 10^{-6})$$

$$= 1 + 1.62 \times 10^{-7}$$

$$= 1.000000162$$

• For  $w_8$  :

$$w_8' = w_8 - \alpha \frac{\partial L}{\partial w_8} = 1 - (0.1)(-1.62 \times 10^{-6})$$

$$\frac{\partial w_8}{\partial w_8} = 1 + 1.62 \times 10^{-7}$$

$$= 1.000000162$$

• For  $w_9$  :

$$w_9' = w_9 - \alpha \frac{\partial L}{\partial w_9} = 1 - 0.1(-1.62 \times 10^{-6})$$

$$\frac{\partial w_9}{\partial w_9} = 1 + 1.62 \times 10^{-7}$$


$$= 1.000000162$$

Comparison :

As all the updated values i.e.  $w_7'$ ,  $w_8'$  and  $w_9'$  are same i.e.  $1.62 \times 10^{-7}$

This is the root cause of symmetry-breaking issue. Symmetry breaking in neural networks refers to situation where weights or neurons in network are initialized or updated where the end up performing same function hindering learning process.





(b)

Given :-

weights  $\Rightarrow w_1 = w_3 = w_5 = -0.5$

$w_2 = w_4 = w_6 = 0.25$

$w_7 = 1, w_8 = -1, w_9 = 0$

biases  $\Rightarrow b_1 = b_2 = b_3 = b_4 = 0.1$

inputs  $\Rightarrow x_1 = 1, x_2 = 1$

Learning Rate ( $\alpha$ ) = 0.1

Using ReLU for hidden layers & sigmoid for output node

As my student ID is 017516785, here the last digit is 5 (odd) hence computing  $w_7, w_1$

- For  $n_1$  :-

$$z_1 = w_1 x_1 + w_2 x_2 + b_1$$

$$= -0.5 \times 1 + 0.25 \times 1 + 0.1$$

$$= -0.5 + 0.25 + 0.1$$

$$= -0.15$$

$$\therefore a_1 = \text{ReLU}(z_1) = \text{ReLU}(-0.15) = 0$$

$$\text{as } \text{ReLU}(x) = \max(0, x)$$

$$\therefore \boxed{a_1 = 0}$$



- For  $n_2 \Rightarrow$

$$\begin{aligned} z_2 &= w_3 x_1 + w_4 x_2 + b_2 \\ &= -0.5 \cdot 1 + 0.25 \cdot 1 + 0.1 \\ &= -0.5 + 0.25 + 0.1 \end{aligned}$$

$$z_2 = -0.15$$

$$a_2 = \text{ReLU}(z_2) = \text{ReLU}(-0.15) = 0$$

- For  $n_3 \Rightarrow$

$$\begin{aligned} z_3 &= w_5 x_1 + w_6 x_2 + b_3 \\ &= -0.5 \cdot 1 + 0.25 \cdot 1 + 0.1 \end{aligned}$$

$$z_3 = -0.15$$

$$a_3 = \text{ReLU}(z_3) = \text{ReLU}(-0.15) = 0$$

- For  $\hat{y}$

$$\begin{aligned} z_4 &= w_7 a_1 + w_8 a_2 + w_9 a_3 + b_4 \\ &= 1 \cdot 0 + (-1) \cdot 0 + 0 \cdot 0 + 0.1 \\ &= 0 + 0 + 0 + 0.1 \end{aligned}$$

$$z_4 = 0.1$$

using sigmoid activation

$$\hat{y} = \sigma(z_4) = \frac{1}{1 + e^{-0.1}} = \frac{1}{1 + 0.9048} = \frac{1}{1.9048}$$

$$\approx 0.52497$$



Loss

$$\begin{aligned}\text{MSE} \Rightarrow L &= \frac{1}{2} (y - \hat{y})^2 = \frac{1}{2} (1 - 0.52497)^2 \\ &= \frac{1}{2} (0.22565) \\ &\approx 0.112826\end{aligned}$$

Backpropagation

Gradient Loss wrt  $\hat{y}$

$$\frac{\partial L}{\partial \hat{y}} = \hat{y} - y = 0.52497 - 1 = -0.47503$$

Gradient Loss wrt  $z_4$

using ~~derivative~~ sigmoid  $\sigma(z) = \frac{1}{1 + e^{-z}}$

$\therefore$  derivative of  $\sigma(z) \Rightarrow$

$$\sigma'(z) = \sigma(z) \cdot (1 - \sigma(z))$$

$$\therefore \sigma'(z_4) = \sigma(z_4) \cdot (1 - \sigma(z_4)) = 0.52497 \cdot (1 - 0.52497)$$

$$\therefore \sigma'(z_4) \approx 0.2494$$

$$\therefore \frac{\partial L}{\partial z_4} = \frac{\partial L}{\partial \hat{y}} \cdot \sigma'(z_4) = -0.4751 \times 0.2494$$

$$\approx -0.1181$$



• For  $w_7$

$$\frac{\partial L}{\partial w_7} = \frac{\partial L}{\partial z_4} \cdot a_1 = -0.1181 \cdot 0 = 0$$

• For  $w_1$

$$\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial z_4} \cdot x_1 = -0.1181 \cdot 1 = -0.1181$$

Hence using learning Rate  $\alpha = 0.1$  we get,

$$w_7' = w \quad \text{Formula} \quad \boxed{w' = w - \alpha \frac{\partial L}{\partial w}}$$

For  $w_7$

$$w_7' = w_7 - \alpha \frac{\partial L}{\partial w_7} = 1 - 0.1 \times 0 = 1$$

• For  $w_1$

$$\begin{aligned} w_1' &= w_1 - \alpha \frac{\partial L}{\partial w_1} = -0.5 - 0.1 \times (-0.1181) \\ &= -0.5 + 0.01181 \\ &= -0.48819 \end{aligned}$$

$\therefore$  For last id odd

$$\underline{w_7' = 1} \quad \text{and} \quad \underline{w_1' = -0.48819}$$



gov98xbx2

February 24, 2025

CODING!

Problem 2-4: We will develop Artificial Neural Networks using MNIST digit data, you can directly download the data using <https://keras.io/api/datasets/mnist/>. The dataset contains 10 classes where each of the image sizes is (28×28). Train for minimum number of epochs = 100; You should split the training data into training and validation sets, and for training the model use these datasets. Test data should be kept separated and used only for evaluation purpose. Your ANN must contain minimum of 2 hidden layers. Apply early stopping criteria based on validation loss with patience 3 and with restoring best weights = true. You may use any regularizers to avoid overfitting.

Problem 2 (8 pts): You should select last two digits of your student ID – meaning that if your student id is 006000104, then you should select 0 and 4 for developing the binary classification model. If both last digit is identical, then select first and last digit. For this task, you must prepare (filter) your data at first to convert the multiclass classification into a binary classification system.

```
[177]: # Importing necessary libraries
import numpy as np
import pandas as pd
import tensorflow as tf
from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import EarlyStopping
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix, classification_report
import warnings

# ignore warnings
warnings.filterwarnings("ignore")

[178]: # to load the MNIST dataset (28x28)
(X_train_full, Y_train_full), (X_test_full, Y_test_full) = mnist.load_data()
```



```
[179]: # to check original shape
print("Original shape of training set:", X_train_full.shape)
print("Original shape of test set:", X_test_full.shape)
```

Original shape of training set: (60000, 28, 28)  
Original shape of test set: (10000, 28, 28)

```
[180]: # to flatten images to 1D vectors [28x28=784]
X_train_full = X_train_full.reshape(X_train_full.shape[0], 784).
    ↳ astype('float32')
X_test_full = X_test_full.reshape(X_test_full.shape[0], 784).astype('float32')
```

```
[181]: # new shape
print("Flattened shape of training set:", X_train_full.shape)
print("Flattened shape of test set:", X_test_full.shape)
```

Flattened shape of training set: (60000, 784)  
Flattened shape of test set: (10000, 784)

```
[182]: # min and max values before normalization
print("Min pixel value (train set):", np.min(X_train_full))
print("Max pixel value (train set):", np.max(X_train_full))
print("Min pixel value (test set):", np.min(X_test_full))
print("Max pixel value (test set):", np.max(X_test_full))
```

Min pixel value (train set): 0.0  
Max pixel value (train set): 255.0  
Min pixel value (test set): 0.0  
Max pixel value (test set): 255.0

```
[183]: # normalize pixel values to the range [0,1]
X_train_full = X_train_full / 255.0
X_test_full = X_test_full / 255.0
```

```
[184]: # min and max values after normalization
print("Min pixel value (train set):", np.min(X_train_full))
print("Max pixel value (train set):", np.max(X_train_full))
print("Min pixel value (test set):", np.min(X_test_full))
print("Max pixel value (test set):", np.max(X_test_full))
```

Min pixel value (train set): 0.0  
Max pixel value (train set): 1.0  
Min pixel value (test set): 0.0  
Max pixel value (test set): 1.0

```
[185]: #selecting 8 and 5 i.e the last two digits from my Student ID = 017516785
selected_digits = [8, 5]
train_filter = np.isin(Y_train_full, selected_digits)
```



```
test_filter = np.isin(Y_test_full, selected_digits)

# filtering
X_train = X_train_full[train_filter]
Y_train = Y_train_full[train_filter]
X_test = X_test_full[test_filter]
Y_test = Y_test_full[test_filter]
```

```
[186]: # new dataset shape
print("Filtered training set shape:", X_train.shape)
print("Filtered test set shape:", X_test.shape)
```

Filtered training set shape: (11272, 784)  
 Filtered test set shape: (1866, 784)

```
[187]: # here convert 5 as 0 , 8 as 1
Y_train = np.where(Y_train == 5, 0, 1)
Y_test = np.where(Y_test == 5, 0, 1)
```

```
[188]: # to check unique values in Y_train and Y_test
print("Unique values in Y_train:", np.unique(Y_train))
print("Unique values in Y_test:", np.unique(Y_test))
```

Unique values in Y\_train: [0 1]  
 Unique values in Y\_test: [0 1]

2a) A. Build an ANN for binary classification. Evaluate your model on the test data. Construct a confusion matrix. Present learning curve (showing training loss and validation loss against number of epochs) and include some examples of your prediction. (3pts)

```
[189]: # Splitting the training set into 80% training and 20% validation
X_train, X_val, Y_train, Y_val = train_test_split(X_train, Y_train, test_size=0.
    ↪2, random_state=42, stratify=Y_train)

# new dataset shapes
print("Training set shape:", X_train.shape)
print("Validation set shape:", X_val.shape)
print("Test set shape:", X_test.shape)
```

Training set shape: (9017, 784)  
 Validation set shape: (2255, 784)  
 Test set shape: (1866, 784)

```
[190]: X_train1=X_train
Y_train1=Y_train
X_test1=X_test
Y_test1=Y_test
X_val1=X_val
```

```
Y_val1=Y_val
```

```
[191]: # ANN model with 2 hidden layer and one output layer
model = Sequential([
    Dense(128, activation='relu', input_shape=(784,)),
    Dense(64, activation='relu'),
    Dense(1, activation='sigmoid')
])

# model compilation
model.compile(loss='binary_crossentropy', metrics=['accuracy'])
```

```
[192]: # early stopping
early_stopping = EarlyStopping(monitor='val_loss', mode='min', patience=3,
    ↪ restore_best_weights=True, verbose=1)
```

```
[193]: model.summary()
```

Model: "sequential\_50"

Layer (type)	Output Shape	
↪ Param #		
dense_194 (Dense)	(None, 128)	↪
↪ 100,480		
dense_195 (Dense)	(None, 64)	↪
↪ 8,256		
dense_196 (Dense)	(None, 1)	↪
↪ 65		

Total params: 108,801 (425.00 KB)

Trainable params: 108,801 (425.00 KB)

Non-trainable params: 0 (0.00 B)

```
[194]: # model train
history = model.fit(
    X_train1, Y_train1,
    validation_data=(X_val1, Y_val1),
```



```
epochs=100,  
batch_size=1000,  
callbacks=[early_stopping],  
verbose=1  
)
```

```
Epoch 1/100  
10/10          3s 118ms/step -  
accuracy: 0.6726 - loss: 0.5494 - val_accuracy: 0.9370 - val_loss: 0.2417  
Epoch 2/100  
10/10          1s 67ms/step -  
accuracy: 0.9442 - loss: 0.2101 - val_accuracy: 0.8905 - val_loss: 0.2731  
Epoch 3/100  
10/10          1s 74ms/step -  
accuracy: 0.9383 - loss: 0.1816 - val_accuracy: 0.9707 - val_loss: 0.1051  
Epoch 4/100  
10/10          1s 86ms/step -  
accuracy: 0.9724 - loss: 0.0966 - val_accuracy: 0.9344 - val_loss: 0.1721  
Epoch 5/100  
10/10          1s 88ms/step -  
accuracy: 0.9521 - loss: 0.1248 - val_accuracy: 0.9627 - val_loss: 0.1049  
Epoch 6/100  
10/10          1s 97ms/step -  
accuracy: 0.9677 - loss: 0.0917 - val_accuracy: 0.9596 - val_loss: 0.1140  
Epoch 7/100  
10/10          1s 67ms/step -  
accuracy: 0.9760 - loss: 0.0748 - val_accuracy: 0.9738 - val_loss: 0.0762  
Epoch 8/100  
10/10          1s 70ms/step -  
accuracy: 0.9796 - loss: 0.0599 - val_accuracy: 0.8652 - val_loss: 0.3313  
Epoch 9/100  
10/10          1s 79ms/step -  
accuracy: 0.9481 - loss: 0.1241 - val_accuracy: 0.9836 - val_loss: 0.0506  
Epoch 10/100  
10/10          1s 59ms/step -  
accuracy: 0.9872 - loss: 0.0404 - val_accuracy: 0.9809 - val_loss: 0.0537  
Epoch 11/100  
10/10          1s 77ms/step -  
accuracy: 0.9834 - loss: 0.0500 - val_accuracy: 0.9854 - val_loss: 0.0456  
Epoch 12/100  
10/10          1s 60ms/step -  
accuracy: 0.9875 - loss: 0.0381 - val_accuracy: 0.9792 - val_loss: 0.0585  
Epoch 13/100  
10/10          1s 50ms/step -  
accuracy: 0.9843 - loss: 0.0447 - val_accuracy: 0.9871 - val_loss: 0.0392  
Epoch 14/100  
10/10          1s 57ms/step -
```

```

accuracy: 0.9912 - loss: 0.0276 - val_accuracy: 0.9814 - val_loss: 0.0522
Epoch 15/100
10/10          1s 59ms/step -
accuracy: 0.9912 - loss: 0.0297 - val_accuracy: 0.9880 - val_loss: 0.0345
Epoch 16/100
10/10          1s 64ms/step -
accuracy: 0.9917 - loss: 0.0261 - val_accuracy: 0.9867 - val_loss: 0.0345
Epoch 17/100
10/10          1s 68ms/step -
accuracy: 0.9933 - loss: 0.0223 - val_accuracy: 0.9885 - val_loss: 0.0326
Epoch 18/100
10/10          1s 118ms/step -
accuracy: 0.9875 - loss: 0.0385 - val_accuracy: 0.9907 - val_loss: 0.0304
Epoch 19/100
10/10          1s 73ms/step -
accuracy: 0.9938 - loss: 0.0188 - val_accuracy: 0.9889 - val_loss: 0.0325
Epoch 20/100
10/10          1s 28ms/step -
accuracy: 0.9905 - loss: 0.0268 - val_accuracy: 0.9902 - val_loss: 0.0304
Epoch 21/100
10/10          1s 27ms/step -
accuracy: 0.9959 - loss: 0.0147 - val_accuracy: 0.9907 - val_loss: 0.0310
Epoch 22/100
10/10          0s 23ms/step -
accuracy: 0.9892 - loss: 0.0293 - val_accuracy: 0.9898 - val_loss: 0.0306
Epoch 23/100
10/10          0s 30ms/step -
accuracy: 0.9964 - loss: 0.0128 - val_accuracy: 0.9477 - val_loss: 0.1541
Epoch 23: early stopping
Restoring model weights from the end of the best epoch: 20.

```

```

[195]: test_loss, test_accuracy = model.evaluate(X_test1, Y_test1)
print(f"Test Loss: {test_loss:.4f}")
print(f"Test Accuracy: {test_accuracy:.4f}")

```

```

59/59          0s 3ms/step -
accuracy: 0.9894 - loss: 0.0283
Test Loss: 0.0252
Test Accuracy: 0.9909

```

### Learning Curves (Training & Validation Loss)

```

[196]: # loss and accuracy data
train_loss = history.history['loss']
val_loss = history.history['val_loss']
train_acc = history.history['accuracy']
val_acc = history.history['val_accuracy']
epochs = range(1, len(train_loss) + 1)

```



```

print("Training Loss:", train_loss)
print("Validation Loss:", val_loss)
print("Training Accuracy:", train_acc)
print("Validation Accuracy:", val_acc)

```

```

Training Loss: [0.4651414752006531, 0.19865594804286957, 0.14709831774234772,
0.09359128773212433, 0.10222889482975006, 0.07810667902231216,
0.06783109158277512, 0.055278949439525604, 0.07918457686901093,
0.04114745184779167, 0.04929564520716667, 0.04039256274700165,
0.03794076293706894, 0.03386365622282028, 0.028625624254345894,
0.027960525825619698, 0.022256115451455116, 0.042767032980918884,
0.018059607595205307, 0.025622360408306122, 0.015745608136057854,
0.026451002806425095, 0.013609780929982662]
Validation Loss: [0.2416839301586151, 0.27312207221984863, 0.10508830100297928,
0.1720648556947708, 0.10490819811820984, 0.11402373015880585,
0.07620465010404587, 0.3313034474849701, 0.05062020942568779,
0.05365185812115669, 0.045562803745269775, 0.058458361774683,
0.03924311324954033, 0.052226755768060684, 0.03446607664227486,
0.034502990543842316, 0.032592177391052246, 0.030359461903572083,
0.032513637095689774, 0.03035561926662922, 0.031041506677865982,
0.03064480423927307, 0.15407925844192505]
Training Accuracy: [0.7671065926551819, 0.945436418056488, 0.9540867209434509,
0.9724963903427124, 0.9627370238304138, 0.9738272428512573, 0.978374183177948,
0.9812576174736023, 0.9693911671638489, 0.9875789880752563, 0.9828102588653564,
0.9866918325424194, 0.9875789880752563, 0.9879117012023926, 0.9920150637626648,
0.9916823506355286, 0.993678629398346, 0.9856936931610107, 0.9942331314086914,
0.9907951354980469, 0.9956748485565186, 0.9911278486251831, 0.996229350566864]
Validation Accuracy: [0.9370288252830505, 0.8904656171798706,
0.9707317352294922, 0.9343680739402771, 0.9627494215965271, 0.9596452116966248,
0.9738359451293945, 0.865188479423523, 0.9835920333862305, 0.980931282043457,
0.9853658676147461, 0.9791574478149414, 0.9871397018432617, 0.9813747406005859,
0.9880266189575195, 0.9866962432861328, 0.9884700775146484, 0.990687370300293,
0.9889135360717773, 0.9902439117431641, 0.990687370300293, 0.9898004531860352,
0.9476718306541443]

```

```


[197]: # two subplots
fig, ax = plt.subplots(1, 2, figsize=(14, 5))

# plot for training & validation Loss
ax[0].plot(epochs, train_loss, 'bo-', label='Training Loss')
ax[0].plot(epochs, val_loss, 'ro-', label='Validation Loss')
ax[0].set_title('Training and Validation Loss')
ax[0].set_xlabel('Epochs')
ax[0].set_ylabel('Loss')
ax[0].legend()

```

```
# plot for training & validation accuracy
ax[1].plot(epochs, train_acc, 'bo-', label='Training Accuracy')
ax[1].plot(epochs, val_acc, 'ro-', label='Validation Accuracy')
ax[1].set_title('Training and Validation Accuracy')
ax[1].set_xlabel('Epochs')
ax[1].set_ylabel('Accuracy')
ax[1].legend()

plt.tight_layout()
plt.show()
```



### Loss Graph:

- The training loss starts high and gradually decreases which mean that the model is learning well.
- The validation loss shows somewhat a similar decreasing trend, telling that the model is generalizing well to unseen data.
- The loss stabilizes after approx 10 epochs which meanss that training beyond this point has little improvement.

**Accuracy Graph:** - Here the training accuracy increases very sharply in the first few epochs and this indicates rapid learning. - Also the validation accuracy closely follows the training accuracy which means that there there is no significant overfitting. - Both training and validation accuracy reach nearly 98 to 99% this suggests that the model performs well on both seen and unseen data.

## 1 Predictions

```
[198]: import random

# random test images
num_samples = 9
indices = random.sample(range(len(X_test1)), num_samples)
sample_images = X_test1[indices].reshape(-1, 28, 28)
sample_labels = Y_test1[indices]
```



```

sample_predictions = (model.predict(X_test1[indices]) > 0.5).astype("int32").
    ↪flatten()


# images with predictions
plt.figure(figsize=(10, 8))
for i in range(num_samples):
    plt.subplot(3, 3, i + 1)
    plt.imshow(sample_images[i], cmap='gray')
    plt.title(f"Actual: {sample_labels[i]} | Predicted:␣
    ↪{sample_predictions[i]}")
    plt.axis("off")

plt.tight_layout()
plt.show()

```

1/1

0s 97ms/step



**Predictions on Randomly Selected Images to test:** - From above image results we can see that the model has correctly identified digits the digits of 8 and 5. - It correctly classifies both 5s and 8s as per their respective labels.

**Interpretation:** - The Model was built with two hidden layers and one output layer. The first hidden layer has 128 units, the second hidden layer has 64 units. Each of these hidden layers uses the ReLU activation function, which is one of the efficient activation functions in neural networks.  
- The model is well-trained and achieves high accuracy in both training and validation.  
- It has a high test accuracy percentage of approx 99% and a low test loss of 0.0252 - No signs of overfitting is observed as the validation loss remains low indicating the model is highly effective in distinguishing between the digits 5 and 8.

[198]:

B. Build ANNs for binary classification using combinations of weight initializers (Normal, He, and Xavier) and activation functions (ReLU, Sigmoid, and tanh). You may use early stopping callback function. Construct confusion matrices and show learning curves for each combination. Create a table, showing combinations, and accuracy. Now, from your experimental results – write a comparative analysis on – impact of different combinations of initializers and activation functions in terms of performance and learning curves. (5 pts)

[199]:

```
from tensorflow.keras import initializers
from sklearn.metrics import confusion_matrix, precision_score, recall_score, f1_score
```

[200]:

```
# to load the MNIST dataset
(X_train, y_train), (X_test, y_test) = mnist.load_data()

# filtering out classes for '8' and '5'
is_eight_or_five_train = (y_train == 8) | (y_train == 5)
is_eight_or_five_test = (y_test == 8) | (y_test == 5)
X_train, y_train = X_train[is_eight_or_five_train], y_train[is_eight_or_five_train]
X_test, y_test = X_test[is_eight_or_five_test], y_test[is_eight_or_five_test]

# to flatten and normalize
X_train = X_train.reshape((-1, 784)).astype('float32') / 255
X_test = X_test.reshape((-1, 784)).astype('float32') / 255

# convert to 1d
y_train = (y_train == 8).astype(int)
y_test = (y_test == 8).astype(int)
```

[201]:

```
def build_and_train_model(initializer, activation_function):
    model = Sequential([
        Dense(128, input_shape=(784,), activation=activation_function, kernel_initializer=initializer),
```



```

        Dense(64, activation=activation_function,
↪kernel_initializer=initializer),
        Dense(32, activation=activation_function,
↪kernel_initializer=initializer),
        Dense(1, activation='sigmoid')
    ])
    model.compile(optimizer='adam', loss='binary_crossentropy',
↪metrics=['accuracy'])
    early_stopping = EarlyStopping(monitor='val_loss', patience=3,
↪restore_best_weights=True)
    history = model.fit(X_train, y_train, validation_split=0.2, epochs=100,
↪callbacks=[early_stopping], verbose=0)
    return model, history

```

[202]: `model.summary()`

Model: "sequential\_50"

Layer (type) ↪Param #	Output Shape	
dense_194 (Dense) ↪100,480	(None, 128)	
dense_195 (Dense) ↪8,256	(None, 64)	
dense_196 (Dense) ↪ 65	(None, 1)	

Total params: 217,604 (850.02 KB)

Trainable params: 108,801 (425.00 KB)

Non-trainable params: 0 (0.00 B)

Optimizer params: 108,803 (425.02 KB)

[203]: `initializers_dict = {`  
`'Normal': initializers.RandomNormal(mean=0.0, stddev=0.05, seed=42),`  
`'He': initializers.HeNormal(seed=42),`

```

    'Xavier': initializers.GlorotUniform(seed=42)
}

activations = ['relu', 'sigmoid', 'tanh']

```

```

[204]: results = []
for name, init in initializers_dict.items():
    for activation in activations:
        model, history = build_and_train_model(init, activation)
        test_loss, test_acc = model.evaluate(X_test, y_test, verbose=0)
        predictions = model.predict(X_test)
        predictions = (predictions > 0.5).astype(int)
        cm = confusion_matrix(y_test, predictions)

        # displaying the confusion matrix
        plt.figure(figsize=(6, 4))
        sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
        plt.title(f'Confusion Matrix: {name} with {activation}')
        plt.show()

        # for plotting the learning curves
        plt.figure(figsize=(12, 4))
        plt.subplot(121)
        plt.plot(history.history['accuracy'], label='Training Accuracy')
        plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
        plt.title(f'Accuracy: {name} with {activation}')
        plt.legend()



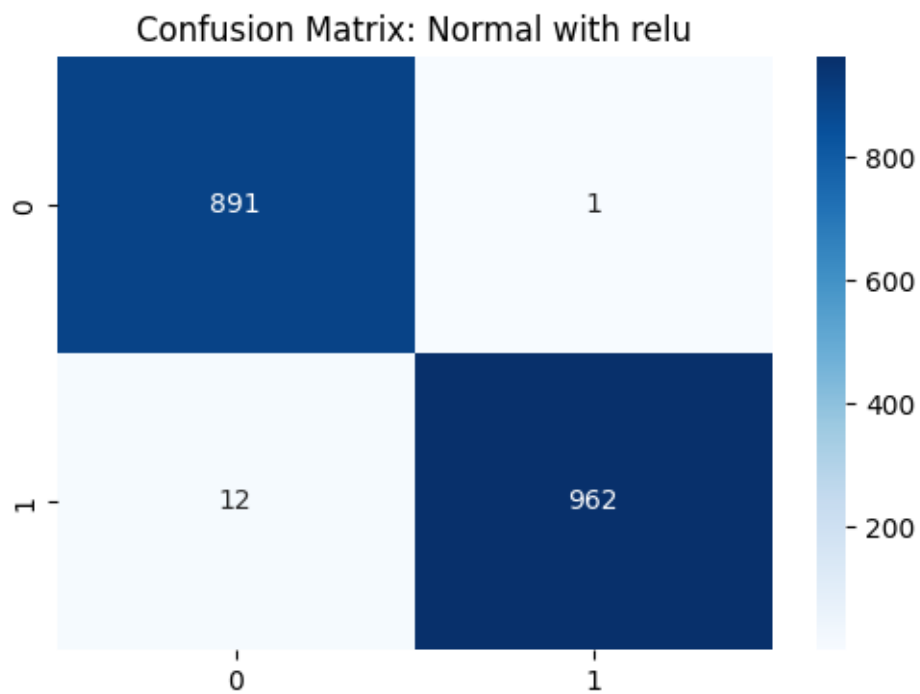
        plt.subplot(122)
        plt.plot(history.history['loss'], label='Training Loss')
        plt.plot(history.history['val_loss'], label='Validation Loss')
        plt.title(f'Loss: {name} with {activation}')
        plt.legend()

        plt.show()

        results.append({
            'Initializer': name,
            'Activation': activation,
            'Accuracy': test_acc,
            'Confusion Matrix': cm
        })

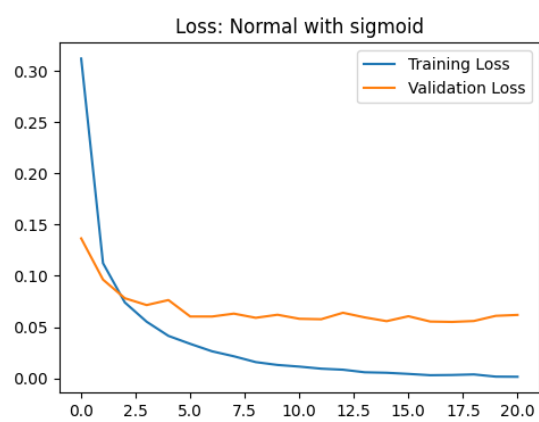

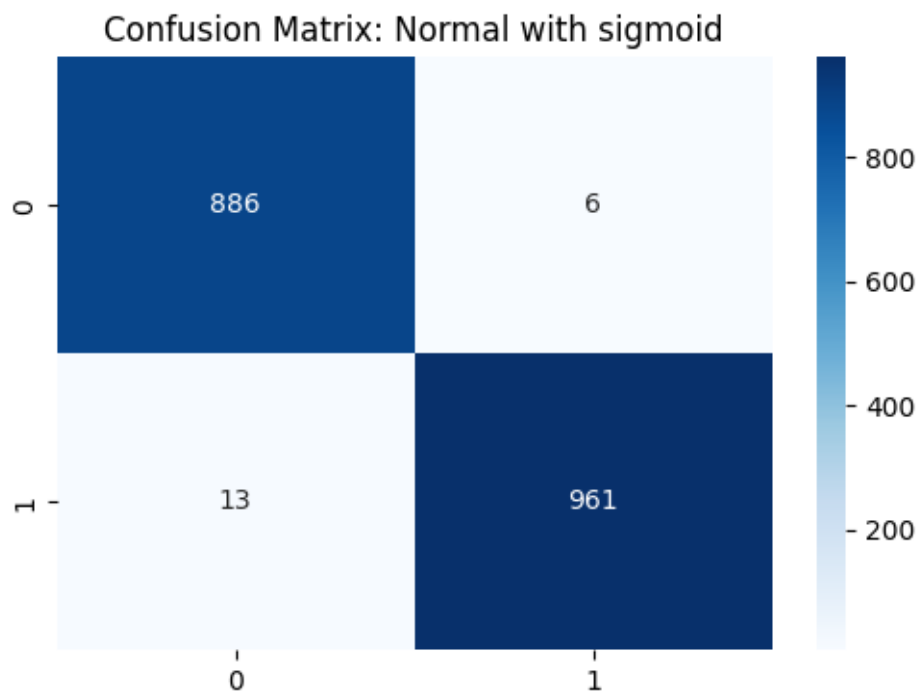
```





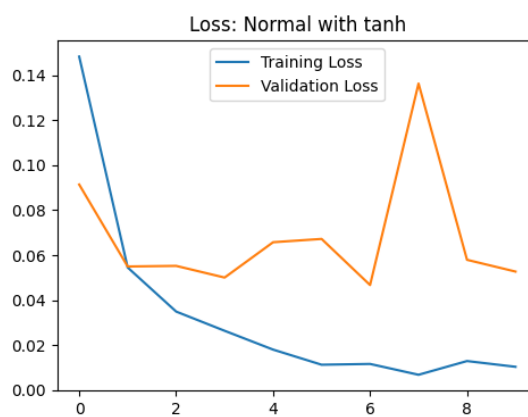

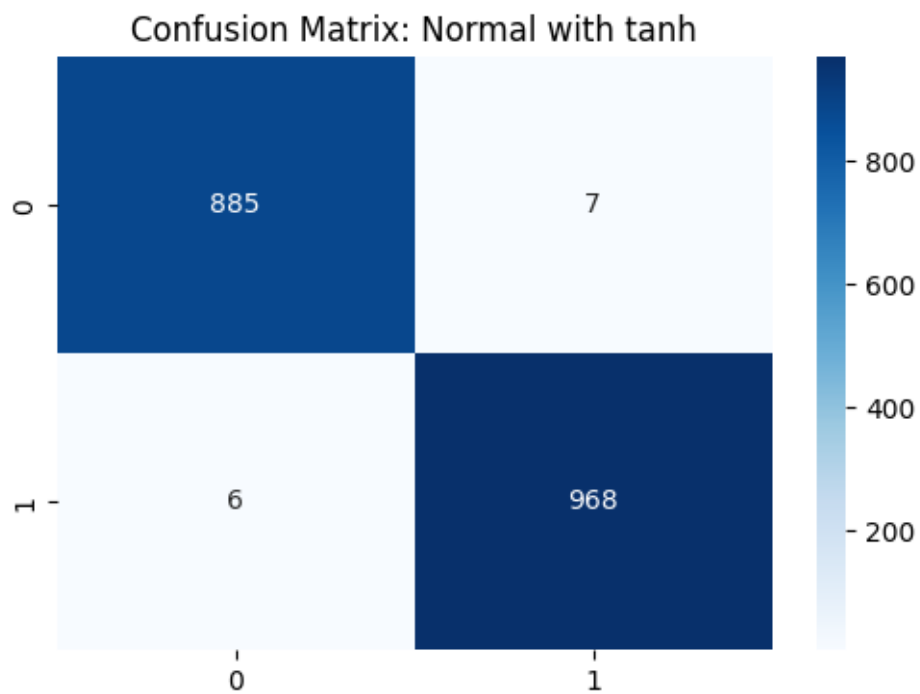
59/59

0s 3ms/step



59/59




0s 3ms/step



59/59




0s 3ms/step








59/59

0s 4ms/step



59/59




0s 3ms/step



59/59

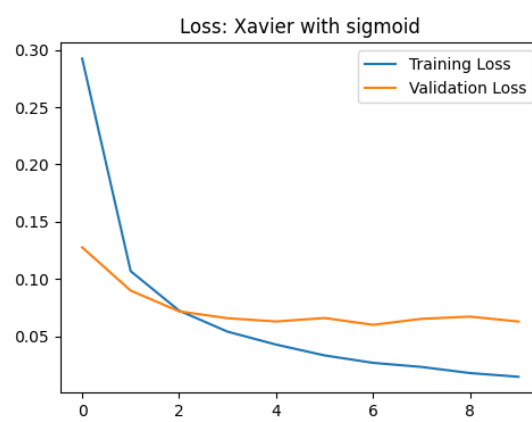


0s 3ms/step







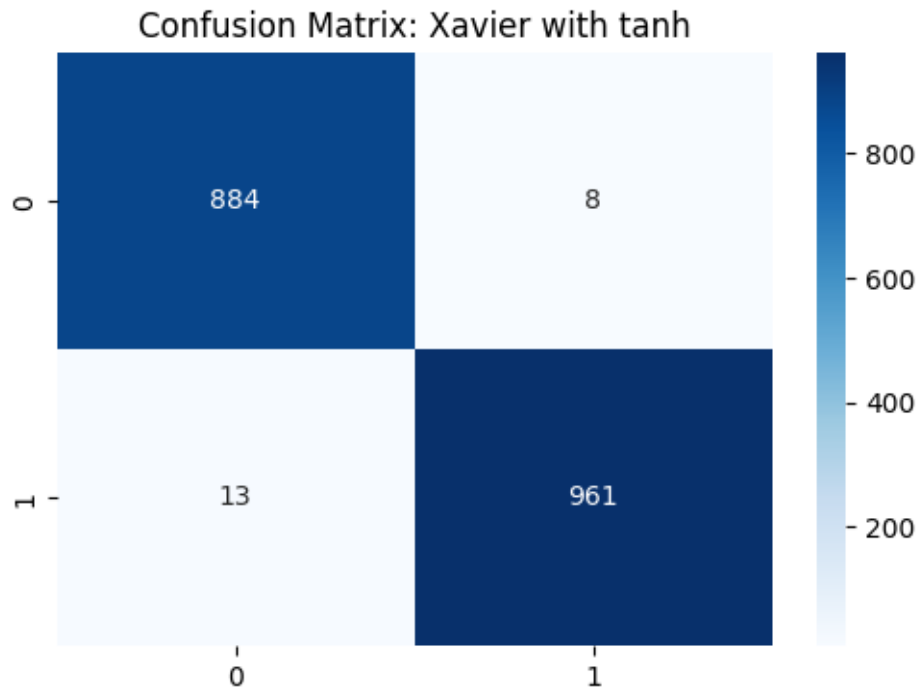
59/59

0s 4ms/step



59/59

0s 3ms/step



```
[205]: precision = precision_score(y_test, predictions)
recall = recall_score(y_test, predictions)
f1 = f1_score(y_test, predictions)

print(f"Precision: {precision}")
print(f"Recall: {recall}")
print(f"F1 Score: {f1}")
```

Precision: 0.9917440660474717  
Recall: 0.9866529774127311



F1 Score: 0.9891919711785898

```
[206]: import pandas as pd
results_df = pd.DataFrame(results)
print(results_df[['Initializer', 'Activation', 'Accuracy']])
```

	Initializer	Activation	Accuracy
0	Normal	relu	0.993033
1	Normal	sigmoid	0.989818
2	Normal	tanh	0.993033
3	He	relu	0.988210
4	He	sigmoid	0.987138
5	He	tanh	0.988746
6	Xavier	relu	0.991961
7	Xavier	sigmoid	0.983923
8	Xavier	tanh	0.988746

**Observations:** - The Normal initializer with ReLU activation had the highest accuracy which is about 99.5% among of the combinations indicating good efficiency. - The He initializer achieved good accuracy results too among which He\_tanh has 99.30% accuracy - ReLU consistently led to quicker and more stable learning, suggesting it's highly effective at managing gradient flow. - Lower performance using Sigmoid points to its vulnerability to gradient issues, requiring careful tuning when used.

[206]:

**1. Normal with Relu Combination:** - Achieves nearly perfect classification with only 1 false positive and 12 false negatives, indicating strong ability in recognizing and distinguishing between the two classes. - Training and validation accuracy both reach high levels above 99%, demonstrating excellent model learning and generalization capabilities without significant overfitting. - Both training and validation loss decrease sharply and stabilize at very low levels, suggesting effective learning with minimal error at the end of training.

**2. Normal with Sigmoid:** - This shows a small increase in classification errors with 6 false positives and 13 false negatives slightly less effective than ReLU but still has high overall accuracy. - Exhibits a stable and consistent accuracy profile after initial learning phases, with both training and validation accuracies leveling above 98%, indicating good model reliability. - The validation loss is notably lower and shows less fluctuation compared to ReLU, pointing to better handling of overfitting and noise within the data.

**3. Normal with Tanh** - Displays increased sensitivity with 7 false positives and 6 false negatives, suggesting potential issues with data extremes that Tanh might not handle as efficiently as ReLU or Sigmoid. - Shows more significant fluctuations in validation accuracy, potentially indicating instability or overfitting issues within the model training process. - Experiences more pronounced spikes in validation loss, especially towards later epochs, which may reflect issues with the model's ability to generalize effectively under Tanh activation.

**4. He with ReLU** - Very high classification accuracy with minimal misclassifications i.e. 9 false positives, 13 false negatives, indicating effective recognition and discrimination between classes. - Has rapid learning with training accuracy quickly reaching a high level, closely mirrored by

validation accuracy which indicates good model generalization without significant overfitting. - The training loss decreases sharply and stabilizes at a very low level, while validation loss also follows but with slight fluctuations, suggesting the model is well-tuned but could possibly benefit from slight adjustments to reduce overfitting.

**5. He with Sigmoid** - Slightly more misclassifications (11 false positives, 13 false negatives) compared to ReLU, indicating a slight reduction in model's discriminative capability. - Both training and validation accuracies show steady increase and plateau at high levels, indicative of effective learning and generalization, albeit slightly below the performance level seen with ReLU. - Training and validation loss both decrease and flatten out, demonstrating stable learning dynamics. The convergence of validation loss close to training loss suggests that the model with sigmoid is robust against overfitting.

**6. He with Tanh** - Has quite a few number of misclassifications to sigmoid (11 false positives, 10 false negatives), showing good performance but slightly less effective than ReLU. - Training accuracy climbs swiftly to nearly 99.5%, while validation accuracy exhibits some fluctuations but remains high, indicating some sensitivity to the model's hyperparameters or the specific characteristics of the validation data. - Both training and validation losses decrease, with training loss achieving a lower level faster. However, the validation loss shows some upward fluctuations later, potentially indicating minor overfitting issues as the model excessively adapts to the training data nuances.

**7. Xavier with ReLU** - This has high classification accuracy with only 5 false positives and 10 false negatives, showing robustness in distinguishing between the two classes. - Training accuracy quickly approaches near-perfect levels, while validation accuracy displays more variability, indicating potential overfitting or sensitivity to validation set nuances. - Both training and validation losses decrease rapidly initially but show notable fluctuations in later epochs, suggesting the need for further parameter tuning or regularization to stabilize learning.

**8. Xavier with Sigmoid** - Observes a minor increase in classification errors with 13 false positives and 17 false negatives compared to the ReLU activation, suggesting a slight decrease in model sensitivity. - Training and validation accuracies rise quickly and plateau at high levels, indicating effective learning and good generalization capabilities of the model. - Displays a smooth and steady reduction in both training and validation losses, indicating consistent learning without abrupt changes, which is beneficial for model stability and performance.

**9. Xavier with Tanh** - Produces a low number of false positives (8) and false negatives (13), demonstrating effective classification capabilities close to those seen with ReLU. - This achieves high training accuracy that stabilizes quickly, although validation accuracy shows some fluctuations, potentially indicating slight overfitting or model sensitivity to validation data specifics. - Shows a decreasing trend in validation loss that levels off, reflecting effective learning adaptation and robustness against overfitting compared to the ReLU and Sigmoid activations.

**Overall Comparison:** - Xavier with ReLU learns quickly but can be unstable, showing ups and downs in its loss. - Sigmoid, however, learns steadily, making it more reliable over time. - Tanh is in the middle, with quick learning but more stable than ReLU. - Sigmoid does the best job in handling new, unseen data because of its consistent performance. - ReLU, while fast, might overfit and not perform as well on new data. - Tanh offers a good balance, performing well on new data without as many ups and downs as ReLU.

[206] :

Problem 3 (2 pts): Build an ANN for multi-class classification considering all the classes (10 classes) in the MNIST digit dataset. Finally, present classification report, including class-wise precision, recall, f1-score, and discuss your result.

```
[207]: from tensorflow.keras.layers import Dense, Flatten
      from tensorflow.keras.utils import to_categorical
      from sklearn.metrics import classification_report, confusion_matrix

[208]: # to load the MNIST dataset (28x28)
      (X_train_full, Y_train_full), (X_test_full, Y_test_full) = mnist.load_data()

      # shape of the data
      print(f"Training data shape: {X_train_full.shape}")
      print(f"Testing data shape: {X_test_full.shape}")
```

Training data shape: (60000, 28, 28)  
Testing data shape: (10000, 28, 28)

```
[209]: # Normalize the data
      X_train = X_train_full / 255.0
      X_test = X_test_full / 255.0

      # to flatten the images from 28x28 to 784
      X_train = X_train.reshape(-1, 28 * 28)
      X_test = X_test.reshape(-1, 28 * 28)

      # One-hot encode the labels
      Y_train = to_categorical(Y_train_full, 10)
      Y_test = to_categorical(Y_test_full, 10)

[210]: # Split in 80% train and 20% validation
      X_train, X_val, Y_train, Y_val = train_test_split(X_train, Y_train, test_size=0.
      ↪2, random_state=42)

      print(f"Training data shape: {X_train.shape}")
      print(f"Validation data shape: {X_val.shape}")
```

Training data shape: (48000, 784)  
Validation data shape: (12000, 784)

```
[211]: # model
      model = Sequential()
      model.add(Dense(128, input_shape=(784,), activation='relu')) # First hidden_
      ↪layer
      model.add(Dense(64, activation='relu')) # Second hidden_
      ↪layer
      model.add(Dense(32, activation='relu')) # Third hidden_
      ↪layer
```



```

model.add(Dense(10, activation='softmax')) # Multi-class
↳ classification so used softmax

model.compile(optimizer='adam', loss='categorical_crossentropy',
↳ metrics=['accuracy'])

```

[212]: model.summary()

Model: "sequential\_60"

Layer (type)	Output Shape	
Param #		
dense_233 (Dense)	(None, 128)	
↳ 100,480		
dense_234 (Dense)	(None, 64)	
↳ 8,256		
dense_235 (Dense)	(None, 32)	
↳ 2,080		
dense_236 (Dense)	(None, 10)	
↳ 330		

Total params: 111,146 (434.16 KB)

Trainable params: 111,146 (434.16 KB)

Non-trainable params: 0 (0.00 B)

[213]: # Early stopping with patience as 3

```

early_stopping = EarlyStopping(monitor='val_loss', patience=3,
↳ restore_best_weights=True)

```

[214]: # train the model with 100 epochs

```

history = model.fit(X_train, Y_train, epochs=100, batch_size=128,
                    validation_data=(X_val, Y_val),
                    callbacks=[early_stopping])

```

Epoch 1/100  
375/375                      5s 10ms/step -

```

accuracy: 0.7757 - loss: 0.7713 - val_accuracy: 0.9405 - val_loss: 0.1993
Epoch 2/100
375/375          6s 16ms/step -
accuracy: 0.9521 - loss: 0.1597 - val_accuracy: 0.9630 - val_loss: 0.1301
Epoch 3/100
375/375          10s 15ms/step -
accuracy: 0.9684 - loss: 0.1061 - val_accuracy: 0.9699 - val_loss: 0.1071
Epoch 4/100
375/375          8s 8ms/step -
accuracy: 0.9762 - loss: 0.0803 - val_accuracy: 0.9707 - val_loss: 0.0993
Epoch 5/100
375/375          6s 11ms/step -
accuracy: 0.9821 - loss: 0.0596 - val_accuracy: 0.9719 - val_loss: 0.0947
Epoch 6/100
375/375          4s 8ms/step -
accuracy: 0.9863 - loss: 0.0482 - val_accuracy: 0.9730 - val_loss: 0.0925
Epoch 7/100
375/375          5s 8ms/step -
accuracy: 0.9890 - loss: 0.0386 - val_accuracy: 0.9693 - val_loss: 0.1049
Epoch 8/100
375/375          9s 18ms/step -
accuracy: 0.9905 - loss: 0.0327 - val_accuracy: 0.9698 - val_loss: 0.1049
Epoch 9/100
375/375          4s 11ms/step -
accuracy: 0.9918 - loss: 0.0286 - val_accuracy: 0.9752 - val_loss: 0.0921
Epoch 10/100
375/375          4s 9ms/step -
accuracy: 0.9945 - loss: 0.0204 - val_accuracy: 0.9755 - val_loss: 0.0982
Epoch 11/100
375/375          3s 8ms/step -
accuracy: 0.9949 - loss: 0.0176 - val_accuracy: 0.9748 - val_loss: 0.1005
Epoch 12/100
375/375          6s 11ms/step -
accuracy: 0.9955 - loss: 0.0166 - val_accuracy: 0.9684 - val_loss: 0.1325

```

```

[215]: # Evaluate the model on the test data
test_loss, test_acc = model.evaluate(X_test, Y_test)
print(f"Test accuracy: {test_acc}")

```

```

313/313          1s 4ms/step -
accuracy: 0.9721 - loss: 0.1033
Test accuracy: 0.9746000170707703

```

```

[216]: from sklearn.metrics import classification_report

Y_pred = model.predict(X_test)

```

```
# to convert predictions from probabilities to class labels
Y_pred_classes = np.argmax(Y_pred, axis=1)

# to convert Y_test to class labels
Y_test_classes = np.argmax(Y_test, axis=1)

# classification report
report_dict = classification_report(Y_test_classes, Y_pred_classes,
    ↪output_dict=True)
print(type(report_dict))
```

```
313/313          2s 5ms/step
<class 'dict'>
```

```
[217]: # to convert dict to a DataFrame
report_df_result = pd.DataFrame(report_dict).transpose()

# report display
display(report_df_result)
```

	precision	recall	f1-score	support
0	0.989754	0.985714	0.987730	980.0000
1	0.989464	0.992952	0.991205	1135.0000
2	0.970221	0.978682	0.974433	1032.0000
3	0.963071	0.981188	0.972045	1010.0000
4	0.963673	0.972505	0.968069	982.0000
5	0.979381	0.958520	0.968839	892.0000
6	0.972079	0.981211	0.976623	958.0000
7	0.979331	0.967899	0.973581	1028.0000
8	0.961382	0.971253	0.966292	974.0000
9	0.976626	0.952428	0.964375	1009.0000
accuracy	0.974600	0.974600	0.974600	0.9746
macro avg	0.974498	0.974235	0.974319	10000.0000
weighted avg	0.974671	0.974600	0.974589	10000.0000

**Observations:** - Since this a multi class classification , we have used softmax activation function in the output layer. - Softmax is suitable because it converts the output scores from the final layer into probabilities - Here we have three hidden layers that use Relu activation funtion to introduce non-linearity into the model, making it capable of learning more complex patterns. - The classification report shows that the model performs well across all digit classes in the MNIST dataset, with an overall accuracy of 97.46%. - Class 1 and Class 7 exhibit the highest precision, indicating that these digits are predicted with the highest accuracy relative to other classes. - Class 5 has the lowest recall, suggesting that it is slightly more challenging for the model to correctly identify all instances of this class compared to others. - The F1-score, which balances precision and recall, is consistently high across all classes, but Class 9 has a slightly lower score, indicating a small trade-off between precision and recall for this class.

```
[217]:
```

[217] :



# cdnvgquol

February 24, 2025

Problem 4 (4 pts): Build ANNs for multi-class classification considering all the classes (10 classes) in the MNIST digit dataset with combinations of batch sizes and learning rates. Consider batch sizes: 4, 16, 32, and 64; and learning rate 0.01, 0.001, 0.0001, and 0.00001. Finally, create a plot of test accuracy vs. ratio of batch size to learning rate, and discuss your findings.

```
[1]: #import necessary libraries
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Flatten, Input
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.callbacks import EarlyStopping

[2]: # Load and preprocess data
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()
train_images = train_images.reshape((60000, 28, 28, 1)).astype('float32') / 255
test_images = test_images.reshape((10000, 28, 28, 1)).astype('float32') / 255
train_labels = to_categorical(train_labels)
test_labels = to_categorical(test_labels)
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>

11490434/11490434                      0s  
0us/step

```
[3]: # model
def create_model(learning_rate):
    model = Sequential([
        Input(shape=(28, 28, 1)),
        Flatten(),
        Dense(32, activation='relu'),
        Dense(16, activation='relu'),
        Dense(10, activation='softmax')
    ])
    model.compile(optimizer=Adam(learning_rate=learning_rate),
```

```

        loss='categorical_crossentropy',
        metrics=['accuracy'])
    return model

```

```
[10]: model.summary()
```

Model: "sequential\_17"

Layer (type) ↳ Param #	Output Shape	
flatten_17 (Flatten) ↳ 0	(None, 784)	↳
dense_51 (Dense) ↳ 25,120	(None, 32)	↳
dense_52 (Dense) ↳ 528	(None, 16)	↳
dense_53 (Dense) ↳ 170	(None, 10)	↳

Total params: 25,818 (100.85 KB)

Trainable params: 25,818 (100.85 KB)

Non-trainable params: 0 (0.00 B)

```

[4]: # earlyStopping
early_stopping = EarlyStopping(
    monitor='val_loss',
    patience=3,
    restore_best_weights=True,
    verbose=1
)

# Batch sizes and learning rates combinations
batch_sizes = [4, 16, 32, 64]
learning_rates = [0.01, 0.001, 0.0001, 0.00001]
results = []

```

```
[5]: # to train models with different batch sizes and learning rates
for batch_size in batch_sizes:
    for learning_rate in learning_rates:
        print(f"Training with batch size {batch_size} and learning rate_{learning_rate}")
        model = create_model(learning_rate=learning_rate)
        model.fit(
            train_images,
            train_labels,
            epochs=100,      #epochs = 100
            batch_size=batch_size,
            validation_split=0.1,
            callbacks=[early_stopping],
            verbose=1
        )
        test_loss, test_accuracy = model.evaluate(test_images, test_labels, verbose=0)
        results.append((batch_size, learning_rate, test_accuracy))
```

Training with batch size 4 and learning rate 0.01

```
Epoch 1/100
13500/13500          21s 1ms/step
- accuracy: 0.8362 - loss: 0.5546 - val_accuracy: 0.9300 - val_loss: 0.2687
Epoch 2/100
13500/13500          19s 1ms/step
- accuracy: 0.9123 - loss: 0.3418 - val_accuracy: 0.9295 - val_loss: 0.2779
Epoch 3/100
13500/13500          20s 1ms/step
- accuracy: 0.9182 - loss: 0.3273 - val_accuracy: 0.9418 - val_loss: 0.2434
Epoch 4/100
13500/13500          21s 2ms/step
- accuracy: 0.9239 - loss: 0.3116 - val_accuracy: 0.9402 - val_loss: 0.2537
Epoch 5/100
13500/13500          20s 1ms/step
- accuracy: 0.9276 - loss: 0.2937 - val_accuracy: 0.9393 - val_loss: 0.2606
Epoch 6/100
13500/13500          22s 2ms/step
- accuracy: 0.9283 - loss: 0.2981 - val_accuracy: 0.9210 - val_loss: 0.3004
Epoch 6: early stopping
```

Restoring model weights from the end of the best epoch: 3.

Training with batch size 4 and learning rate 0.001

```
Epoch 1/100
13500/13500          21s 1ms/step
- accuracy: 0.8597 - loss: 0.4568 - val_accuracy: 0.9553 - val_loss: 0.1553
Epoch 2/100
13500/13500          21s 2ms/step
- accuracy: 0.9531 - loss: 0.1598 - val_accuracy: 0.9667 - val_loss: 0.1225
```

Epoch 3/100  
13500/13500 40s 1ms/step  
- accuracy: 0.9625 - loss: 0.1223 - val\_accuracy: 0.9678 - val\_loss: 0.1087

Epoch 4/100  
13500/13500 21s 1ms/step  
- accuracy: 0.9697 - loss: 0.1005 - val\_accuracy: 0.9635 - val\_loss: 0.1352

Epoch 5/100  
13500/13500 21s 1ms/step  
- accuracy: 0.9722 - loss: 0.0899 - val\_accuracy: 0.9675 - val\_loss: 0.1199

Epoch 6/100  
13500/13500 21s 2ms/step  
- accuracy: 0.9750 - loss: 0.0809 - val\_accuracy: 0.9645 - val\_loss: 0.1304

Epoch 6: early stopping  
Restoring model weights from the end of the best epoch: 3.  
Training with batch size 4 and learning rate 0.0001

Epoch 1/100  
13500/13500 21s 2ms/step  
- accuracy: 0.6951 - loss: 1.0041 - val\_accuracy: 0.9285 - val\_loss: 0.2665

Epoch 2/100  
13500/13500 40s 1ms/step  
- accuracy: 0.9153 - loss: 0.3030 - val\_accuracy: 0.9358 - val\_loss: 0.2223

Epoch 3/100  
13500/13500 22s 2ms/step  
- accuracy: 0.9300 - loss: 0.2441 - val\_accuracy: 0.9495 - val\_loss: 0.1841

Epoch 4/100  
13500/13500 21s 2ms/step  
- accuracy: 0.9400 - loss: 0.2090 - val\_accuracy: 0.9567 - val\_loss: 0.1647

Epoch 5/100  
13500/13500 41s 2ms/step  
- accuracy: 0.9463 - loss: 0.1833 - val\_accuracy: 0.9593 - val\_loss: 0.1525

Epoch 6/100  
13500/13500 20s 2ms/step  
- accuracy: 0.9528 - loss: 0.1643 - val\_accuracy: 0.9597 - val\_loss: 0.1462

Epoch 7/100  
13500/13500 21s 2ms/step  
- accuracy: 0.9555 - loss: 0.1521 - val\_accuracy: 0.9627 - val\_loss: 0.1383

Epoch 8/100  
13500/13500 43s 2ms/step  
- accuracy: 0.9587 - loss: 0.1443 - val\_accuracy: 0.9630 - val\_loss: 0.1320

Epoch 9/100  
13500/13500 40s 2ms/step  
- accuracy: 0.9605 - loss: 0.1307 - val\_accuracy: 0.9652 - val\_loss: 0.1317

Epoch 10/100  
13500/13500 21s 2ms/step  
- accuracy: 0.9639 - loss: 0.1238 - val\_accuracy: 0.9648 - val\_loss: 0.1277

Epoch 11/100  
13500/13500 40s 2ms/step  
- accuracy: 0.9656 - loss: 0.1186 - val\_accuracy: 0.9655 - val\_loss: 0.1242

Epoch 12/100  
13500/13500 21s 2ms/step  
- accuracy: 0.9678 - loss: 0.1098 - val\_accuracy: 0.9672 - val\_loss: 0.1191

Epoch 13/100  
13500/13500 21s 2ms/step  
- accuracy: 0.9677 - loss: 0.1063 - val\_accuracy: 0.9668 - val\_loss: 0.1191

Epoch 14/100  
13500/13500 41s 2ms/step  
- accuracy: 0.9708 - loss: 0.1010 - val\_accuracy: 0.9653 - val\_loss: 0.1210

Epoch 15/100  
13500/13500 21s 2ms/step  
- accuracy: 0.9719 - loss: 0.0949 - val\_accuracy: 0.9682 - val\_loss: 0.1141

Epoch 16/100  
13500/13500 21s 2ms/step  
- accuracy: 0.9741 - loss: 0.0899 - val\_accuracy: 0.9697 - val\_loss: 0.1157

Epoch 17/100  
13500/13500 41s 2ms/step  
- accuracy: 0.9750 - loss: 0.0868 - val\_accuracy: 0.9703 - val\_loss: 0.1118

Epoch 18/100  
13500/13500 41s 2ms/step  
- accuracy: 0.9765 - loss: 0.0800 - val\_accuracy: 0.9708 - val\_loss: 0.1127

Epoch 19/100  
13500/13500 41s 2ms/step  
- accuracy: 0.9758 - loss: 0.0829 - val\_accuracy: 0.9713 - val\_loss: 0.1122

Epoch 20/100  
13500/13500 41s 2ms/step  
- accuracy: 0.9779 - loss: 0.0762 - val\_accuracy: 0.9700 - val\_loss: 0.1166

Epoch 20: early stopping  
Restoring model weights from the end of the best epoch: 17.  
Training with batch size 4 and learning rate 1e-05

Epoch 1/100  
13500/13500 23s 2ms/step  
- accuracy: 0.3344 - loss: 1.9746 - val\_accuracy: 0.7358 - val\_loss: 1.0314

Epoch 2/100  
13500/13500 40s 2ms/step  
- accuracy: 0.7638 - loss: 0.9308 - val\_accuracy: 0.8717 - val\_loss: 0.5906

Epoch 3/100  
13500/13500 41s 2ms/step  
- accuracy: 0.8535 - loss: 0.6101 - val\_accuracy: 0.8998 - val\_loss: 0.4452

Epoch 4/100  
13500/13500 41s 2ms/step  
- accuracy: 0.8771 - loss: 0.4883 - val\_accuracy: 0.9113 - val\_loss: 0.3757

Epoch 5/100  
13500/13500 40s 2ms/step  
- accuracy: 0.8903 - loss: 0.4262 - val\_accuracy: 0.9167 - val\_loss: 0.3356

Epoch 6/100  
13500/13500 41s 2ms/step  
- accuracy: 0.8959 - loss: 0.3883 - val\_accuracy: 0.9228 - val\_loss: 0.3083



Epoch 7/100  
13500/13500 21s 2ms/step  
- accuracy: 0.9042 - loss: 0.3603 - val\_accuracy: 0.9257 - val\_loss: 0.2902

Epoch 8/100  
13500/13500 41s 2ms/step  
- accuracy: 0.9073 - loss: 0.3476 - val\_accuracy: 0.9280 - val\_loss: 0.2758

Epoch 9/100  
13500/13500 42s 2ms/step  
- accuracy: 0.9124 - loss: 0.3221 - val\_accuracy: 0.9292 - val\_loss: 0.2650

Epoch 10/100  
13500/13500 41s 2ms/step  
- accuracy: 0.9142 - loss: 0.3123 - val\_accuracy: 0.9317 - val\_loss: 0.2557

Epoch 11/100  
13500/13500 21s 2ms/step  
- accuracy: 0.9172 - loss: 0.3057 - val\_accuracy: 0.9330 - val\_loss: 0.2481

Epoch 12/100  
13500/13500 21s 2ms/step  
- accuracy: 0.9193 - loss: 0.2913 - val\_accuracy: 0.9350 - val\_loss: 0.2425

Epoch 13/100  
13500/13500 21s 2ms/step  
- accuracy: 0.9207 - loss: 0.2871 - val\_accuracy: 0.9355 - val\_loss: 0.2365

Epoch 14/100  
13500/13500 41s 2ms/step  
- accuracy: 0.9224 - loss: 0.2800 - val\_accuracy: 0.9370 - val\_loss: 0.2319

Epoch 15/100  
13500/13500 22s 2ms/step  
- accuracy: 0.9249 - loss: 0.2735 - val\_accuracy: 0.9377 - val\_loss: 0.2259

Epoch 16/100  
13500/13500 22s 2ms/step  
- accuracy: 0.9241 - loss: 0.2678 - val\_accuracy: 0.9390 - val\_loss: 0.2218

Epoch 17/100  
13500/13500 23s 2ms/step  
- accuracy: 0.9263 - loss: 0.2587 - val\_accuracy: 0.9383 - val\_loss: 0.2187

Epoch 18/100  
13500/13500 22s 2ms/step  
- accuracy: 0.9281 - loss: 0.2592 - val\_accuracy: 0.9400 - val\_loss: 0.2151

Epoch 19/100  
13500/13500 41s 2ms/step  
- accuracy: 0.9281 - loss: 0.2510 - val\_accuracy: 0.9415 - val\_loss: 0.2120

Epoch 20/100  
13500/13500 22s 2ms/step  
- accuracy: 0.9305 - loss: 0.2475 - val\_accuracy: 0.9425 - val\_loss: 0.2081

Epoch 21/100  
13500/13500 42s 2ms/step  
- accuracy: 0.9315 - loss: 0.2406 - val\_accuracy: 0.9432 - val\_loss: 0.2055

Epoch 22/100  
13500/13500 39s 2ms/step  
- accuracy: 0.9315 - loss: 0.2434 - val\_accuracy: 0.9448 - val\_loss: 0.2029

Epoch 23/100  
13500/13500 42s 2ms/step  
- accuracy: 0.9309 - loss: 0.2431 - val\_accuracy: 0.9438 - val\_loss: 0.1996

Epoch 24/100  
13500/13500 41s 2ms/step  
- accuracy: 0.9334 - loss: 0.2349 - val\_accuracy: 0.9463 - val\_loss: 0.1970

Epoch 25/100  
13500/13500 21s 2ms/step  
- accuracy: 0.9369 - loss: 0.2275 - val\_accuracy: 0.9468 - val\_loss: 0.1941

Epoch 26/100  
13500/13500 42s 2ms/step  
- accuracy: 0.9361 - loss: 0.2268 - val\_accuracy: 0.9483 - val\_loss: 0.1919

Epoch 27/100  
13500/13500 22s 2ms/step  
- accuracy: 0.9365 - loss: 0.2250 - val\_accuracy: 0.9482 - val\_loss: 0.1894

Epoch 28/100  
13500/13500 41s 2ms/step  
- accuracy: 0.9373 - loss: 0.2210 - val\_accuracy: 0.9480 - val\_loss: 0.1875

Epoch 29/100  
13500/13500 41s 2ms/step  
- accuracy: 0.9390 - loss: 0.2190 - val\_accuracy: 0.9485 - val\_loss: 0.1853

Epoch 30/100  
13500/13500 21s 2ms/step  
- accuracy: 0.9390 - loss: 0.2165 - val\_accuracy: 0.9495 - val\_loss: 0.1838

Epoch 31/100  
13500/13500 41s 2ms/step  
- accuracy: 0.9387 - loss: 0.2176 - val\_accuracy: 0.9502 - val\_loss: 0.1842

Epoch 32/100  
13500/13500 22s 2ms/step  
- accuracy: 0.9408 - loss: 0.2093 - val\_accuracy: 0.9502 - val\_loss: 0.1798

Epoch 33/100  
13500/13500 40s 2ms/step  
- accuracy: 0.9407 - loss: 0.2077 - val\_accuracy: 0.9505 - val\_loss: 0.1791

Epoch 34/100  
13500/13500 21s 2ms/step  
- accuracy: 0.9422 - loss: 0.2045 - val\_accuracy: 0.9513 - val\_loss: 0.1765

Epoch 35/100  
13500/13500 22s 2ms/step  
- accuracy: 0.9427 - loss: 0.2014 - val\_accuracy: 0.9525 - val\_loss: 0.1750

Epoch 36/100  
13500/13500 22s 2ms/step  
- accuracy: 0.9436 - loss: 0.1982 - val\_accuracy: 0.9517 - val\_loss: 0.1737

Epoch 37/100  
13500/13500 20s 2ms/step  
- accuracy: 0.9433 - loss: 0.2015 - val\_accuracy: 0.9528 - val\_loss: 0.1722

Epoch 38/100  
13500/13500 22s 2ms/step  
- accuracy: 0.9450 - loss: 0.1968 - val\_accuracy: 0.9543 - val\_loss: 0.1710

Epoch 39/100  
13500/13500 40s 2ms/step  
- accuracy: 0.9457 - loss: 0.1920 - val\_accuracy: 0.9533 - val\_loss: 0.1695

Epoch 40/100  
13500/13500 21s 2ms/step  
- accuracy: 0.9430 - loss: 0.1996 - val\_accuracy: 0.9540 - val\_loss: 0.1682

Epoch 41/100  
13500/13500 22s 2ms/step  
- accuracy: 0.9470 - loss: 0.1882 - val\_accuracy: 0.9547 - val\_loss: 0.1676

Epoch 42/100  
13500/13500 21s 2ms/step  
- accuracy: 0.9474 - loss: 0.1890 - val\_accuracy: 0.9543 - val\_loss: 0.1659

Epoch 43/100  
13500/13500 22s 2ms/step  
- accuracy: 0.9472 - loss: 0.1859 - val\_accuracy: 0.9548 - val\_loss: 0.1651

Epoch 44/100  
13500/13500 22s 2ms/step  
- accuracy: 0.9474 - loss: 0.1863 - val\_accuracy: 0.9552 - val\_loss: 0.1637

Epoch 45/100  
13500/13500 41s 2ms/step  
- accuracy: 0.9492 - loss: 0.1804 - val\_accuracy: 0.9545 - val\_loss: 0.1628

Epoch 46/100  
13500/13500 41s 2ms/step  
- accuracy: 0.9486 - loss: 0.1793 - val\_accuracy: 0.9565 - val\_loss: 0.1616

Epoch 47/100  
13500/13500 40s 2ms/step  
- accuracy: 0.9500 - loss: 0.1767 - val\_accuracy: 0.9562 - val\_loss: 0.1609

Epoch 48/100  
13500/13500 40s 2ms/step  
- accuracy: 0.9488 - loss: 0.1803 - val\_accuracy: 0.9565 - val\_loss: 0.1609

Epoch 49/100  
13500/13500 22s 2ms/step  
- accuracy: 0.9500 - loss: 0.1771 - val\_accuracy: 0.9580 - val\_loss: 0.1592

Epoch 50/100  
13500/13500 40s 2ms/step  
- accuracy: 0.9502 - loss: 0.1757 - val\_accuracy: 0.9565 - val\_loss: 0.1583

Epoch 51/100  
13500/13500 42s 2ms/step  
- accuracy: 0.9524 - loss: 0.1728 - val\_accuracy: 0.9573 - val\_loss: 0.1568

Epoch 52/100  
13500/13500 42s 2ms/step  
- accuracy: 0.9503 - loss: 0.1734 - val\_accuracy: 0.9582 - val\_loss: 0.1563

Epoch 53/100  
13500/13500 23s 2ms/step  
- accuracy: 0.9511 - loss: 0.1711 - val\_accuracy: 0.9580 - val\_loss: 0.1552

Epoch 54/100  
13500/13500 41s 2ms/step  
- accuracy: 0.9540 - loss: 0.1634 - val\_accuracy: 0.9585 - val\_loss: 0.1545

Epoch 55/100  
13500/13500 41s 2ms/step  
- accuracy: 0.9507 - loss: 0.1742 - val\_accuracy: 0.9598 - val\_loss: 0.1535

Epoch 56/100  
13500/13500 41s 2ms/step  
- accuracy: 0.9518 - loss: 0.1714 - val\_accuracy: 0.9588 - val\_loss: 0.1529

Epoch 57/100  
13500/13500 22s 2ms/step  
- accuracy: 0.9529 - loss: 0.1631 - val\_accuracy: 0.9598 - val\_loss: 0.1523

Epoch 58/100  
13500/13500 40s 2ms/step  
- accuracy: 0.9532 - loss: 0.1649 - val\_accuracy: 0.9600 - val\_loss: 0.1517

Epoch 59/100  
13500/13500 21s 2ms/step  
- accuracy: 0.9530 - loss: 0.1606 - val\_accuracy: 0.9597 - val\_loss: 0.1505

Epoch 60/100  
13500/13500 42s 2ms/step  
- accuracy: 0.9544 - loss: 0.1606 - val\_accuracy: 0.9612 - val\_loss: 0.1492

Epoch 61/100  
13500/13500 21s 2ms/step  
- accuracy: 0.9547 - loss: 0.1577 - val\_accuracy: 0.9605 - val\_loss: 0.1491

Epoch 62/100  
13500/13500 21s 2ms/step  
- accuracy: 0.9549 - loss: 0.1565 - val\_accuracy: 0.9612 - val\_loss: 0.1477

Epoch 63/100  
13500/13500 41s 2ms/step  
- accuracy: 0.9561 - loss: 0.1575 - val\_accuracy: 0.9607 - val\_loss: 0.1474

Epoch 64/100  
13500/13500 22s 2ms/step  
- accuracy: 0.9552 - loss: 0.1588 - val\_accuracy: 0.9605 - val\_loss: 0.1473

Epoch 65/100  
13500/13500 21s 2ms/step  
- accuracy: 0.9573 - loss: 0.1560 - val\_accuracy: 0.9610 - val\_loss: 0.1464

Epoch 66/100  
13500/13500 21s 2ms/step  
- accuracy: 0.9575 - loss: 0.1504 - val\_accuracy: 0.9617 - val\_loss: 0.1460

Epoch 67/100  
13500/13500 43s 2ms/step  
- accuracy: 0.9581 - loss: 0.1480 - val\_accuracy: 0.9603 - val\_loss: 0.1451

Epoch 68/100  
13500/13500 40s 2ms/step  
- accuracy: 0.9569 - loss: 0.1541 - val\_accuracy: 0.9607 - val\_loss: 0.1447

Epoch 69/100  
13500/13500 42s 2ms/step  
- accuracy: 0.9582 - loss: 0.1462 - val\_accuracy: 0.9612 - val\_loss: 0.1439

Epoch 70/100  
13500/13500 22s 2ms/step  
- accuracy: 0.9586 - loss: 0.1479 - val\_accuracy: 0.9613 - val\_loss: 0.1437

Epoch 71/100  
13500/13500 40s 2ms/step  
- accuracy: 0.9589 - loss: 0.1427 - val\_accuracy: 0.9617 - val\_loss: 0.1430

Epoch 72/100  
13500/13500 21s 2ms/step  
- accuracy: 0.9583 - loss: 0.1488 - val\_accuracy: 0.9617 - val\_loss: 0.1427

Epoch 73/100  
13500/13500 22s 2ms/step  
- accuracy: 0.9605 - loss: 0.1409 - val\_accuracy: 0.9622 - val\_loss: 0.1420

Epoch 74/100  
13500/13500 22s 2ms/step  
- accuracy: 0.9598 - loss: 0.1439 - val\_accuracy: 0.9620 - val\_loss: 0.1423

Epoch 75/100  
13500/13500 41s 2ms/step  
- accuracy: 0.9590 - loss: 0.1437 - val\_accuracy: 0.9620 - val\_loss: 0.1413

Epoch 76/100  
13500/13500 40s 2ms/step  
- accuracy: 0.9593 - loss: 0.1426 - val\_accuracy: 0.9618 - val\_loss: 0.1412

Epoch 77/100  
13500/13500 41s 2ms/step  
- accuracy: 0.9594 - loss: 0.1435 - val\_accuracy: 0.9627 - val\_loss: 0.1404

Epoch 78/100  
13500/13500 22s 2ms/step  
- accuracy: 0.9622 - loss: 0.1365 - val\_accuracy: 0.9617 - val\_loss: 0.1404

Epoch 79/100  
13500/13500 22s 2ms/step  
- accuracy: 0.9608 - loss: 0.1401 - val\_accuracy: 0.9622 - val\_loss: 0.1398

Epoch 80/100  
13500/13500 22s 2ms/step  
- accuracy: 0.9617 - loss: 0.1367 - val\_accuracy: 0.9617 - val\_loss: 0.1391

Epoch 81/100  
13500/13500 41s 2ms/step  
- accuracy: 0.9608 - loss: 0.1398 - val\_accuracy: 0.9612 - val\_loss: 0.1394

Epoch 82/100  
13500/13500 41s 2ms/step  
- accuracy: 0.9608 - loss: 0.1358 - val\_accuracy: 0.9620 - val\_loss: 0.1386

Epoch 83/100  
13500/13500 21s 2ms/step  
- accuracy: 0.9612 - loss: 0.1367 - val\_accuracy: 0.9625 - val\_loss: 0.1383

Epoch 84/100  
13500/13500 42s 2ms/step  
- accuracy: 0.9622 - loss: 0.1310 - val\_accuracy: 0.9625 - val\_loss: 0.1384

Epoch 85/100  
13500/13500 41s 2ms/step  
- accuracy: 0.9626 - loss: 0.1319 - val\_accuracy: 0.9630 - val\_loss: 0.1370

Epoch 86/100  
13500/13500 40s 2ms/step  
- accuracy: 0.9619 - loss: 0.1353 - val\_accuracy: 0.9625 - val\_loss: 0.1369



Epoch 87/100  
13500/13500 41s 2ms/step  
- accuracy: 0.9633 - loss: 0.1323 - val\_accuracy: 0.9628 - val\_loss: 0.1363  
Epoch 88/100  
13500/13500 22s 2ms/step  
- accuracy: 0.9620 - loss: 0.1322 - val\_accuracy: 0.9627 - val\_loss: 0.1356  
Epoch 89/100  
13500/13500 41s 2ms/step  
- accuracy: 0.9640 - loss: 0.1303 - val\_accuracy: 0.9623 - val\_loss: 0.1356  
Epoch 90/100  
13500/13500 22s 2ms/step  
- accuracy: 0.9629 - loss: 0.1308 - val\_accuracy: 0.9637 - val\_loss: 0.1349  
Epoch 91/100  
13500/13500 22s 2ms/step  
- accuracy: 0.9627 - loss: 0.1295 - val\_accuracy: 0.9640 - val\_loss: 0.1353  
Epoch 92/100  
13500/13500 22s 2ms/step  
- accuracy: 0.9653 - loss: 0.1236 - val\_accuracy: 0.9635 - val\_loss: 0.1343  
Epoch 93/100  
13500/13500 23s 2ms/step  
- accuracy: 0.9639 - loss: 0.1269 - val\_accuracy: 0.9627 - val\_loss: 0.1348  
Epoch 94/100  
13500/13500 22s 2ms/step  
- accuracy: 0.9633 - loss: 0.1279 - val\_accuracy: 0.9640 - val\_loss: 0.1343  
Epoch 95/100  
13500/13500 41s 2ms/step  
- accuracy: 0.9647 - loss: 0.1247 - val\_accuracy: 0.9638 - val\_loss: 0.1332  
Epoch 96/100  
13500/13500 22s 2ms/step  
- accuracy: 0.9654 - loss: 0.1225 - val\_accuracy: 0.9647 - val\_loss: 0.1334  
Epoch 97/100  
13500/13500 41s 2ms/step  
- accuracy: 0.9644 - loss: 0.1261 - val\_accuracy: 0.9638 - val\_loss: 0.1331  
Epoch 98/100  
13500/13500 22s 2ms/step  
- accuracy: 0.9642 - loss: 0.1267 - val\_accuracy: 0.9638 - val\_loss: 0.1330  
Epoch 99/100  
13500/13500 40s 2ms/step  
- accuracy: 0.9652 - loss: 0.1224 - val\_accuracy: 0.9635 - val\_loss: 0.1323  
Epoch 100/100  
13500/13500 42s 2ms/step  
- accuracy: 0.9660 - loss: 0.1222 - val\_accuracy: 0.9640 - val\_loss: 0.1322  
Restoring model weights from the end of the best epoch: 100.  
Training with batch size 16 and learning rate 0.01  
Epoch 1/100  
3375/3375 7s 2ms/step -  
accuracy: 0.8560 - loss: 0.4736 - val\_accuracy: 0.9420 - val\_loss: 0.2078  
Epoch 2/100

```

3375/3375          6s 2ms/step -
accuracy: 0.9285 - loss: 0.2504 - val_accuracy: 0.9467 - val_loss: 0.1920
Epoch 3/100
3375/3375          10s 2ms/step -
accuracy: 0.9375 - loss: 0.2260 - val_accuracy: 0.9487 - val_loss: 0.2034
Epoch 4/100
3375/3375          6s 2ms/step -
accuracy: 0.9423 - loss: 0.2101 - val_accuracy: 0.9495 - val_loss: 0.1917
Epoch 5/100
3375/3375          5s 2ms/step -
accuracy: 0.9446 - loss: 0.2098 - val_accuracy: 0.9562 - val_loss: 0.1916
Epoch 6/100
3375/3375          10s 2ms/step -
accuracy: 0.9486 - loss: 0.1998 - val_accuracy: 0.9567 - val_loss: 0.1709
Epoch 7/100
3375/3375          5s 2ms/step -
accuracy: 0.9519 - loss: 0.1755 - val_accuracy: 0.9505 - val_loss: 0.2025
Epoch 8/100
3375/3375          6s 2ms/step -
accuracy: 0.9521 - loss: 0.1819 - val_accuracy: 0.9600 - val_loss: 0.1771
Epoch 9/100
3375/3375          10s 2ms/step -
accuracy: 0.9515 - loss: 0.1828 - val_accuracy: 0.9538 - val_loss: 0.1932
Epoch 9: early stopping
Restoring model weights from the end of the best epoch: 6.
Training with batch size 16 and learning rate 0.001
Epoch 1/100
3375/3375          7s 2ms/step -
accuracy: 0.7944 - loss: 0.6861 - val_accuracy: 0.9552 - val_loss: 0.1680
Epoch 2/100
3375/3375          10s 2ms/step -
accuracy: 0.9397 - loss: 0.2025 - val_accuracy: 0.9588 - val_loss: 0.1385
Epoch 3/100
3375/3375          10s 2ms/step -
accuracy: 0.9540 - loss: 0.1533 - val_accuracy: 0.9647 - val_loss: 0.1232
Epoch 4/100
3375/3375          6s 2ms/step -
accuracy: 0.9627 - loss: 0.1220 - val_accuracy: 0.9680 - val_loss: 0.1139
Epoch 5/100
3375/3375          10s 2ms/step -
accuracy: 0.9680 - loss: 0.1056 - val_accuracy: 0.9673 - val_loss: 0.1120
Epoch 6/100
3375/3375          10s 2ms/step -
accuracy: 0.9712 - loss: 0.0926 - val_accuracy: 0.9730 - val_loss: 0.1033
Epoch 7/100
3375/3375          6s 2ms/step -
accuracy: 0.9739 - loss: 0.0850 - val_accuracy: 0.9687 - val_loss: 0.1048
Epoch 8/100

```

```

3375/3375          6s 2ms/step -
accuracy: 0.9774 - loss: 0.0722 - val_accuracy: 0.9723 - val_loss: 0.1005
Epoch 9/100
3375/3375          6s 2ms/step -
accuracy: 0.9784 - loss: 0.0683 - val_accuracy: 0.9705 - val_loss: 0.1116
Epoch 10/100
3375/3375          10s 2ms/step -
accuracy: 0.9798 - loss: 0.0620 - val_accuracy: 0.9698 - val_loss: 0.1012
Epoch 11/100
3375/3375          11s 2ms/step -
accuracy: 0.9810 - loss: 0.0562 - val_accuracy: 0.9708 - val_loss: 0.1100
Epoch 11: early stopping
Restoring model weights from the end of the best epoch: 8.
Training with batch size 16 and learning rate 0.0001
Epoch 1/100
3375/3375          7s 2ms/step -
accuracy: 0.5906 - loss: 1.3346 - val_accuracy: 0.9173 - val_loss: 0.3384
Epoch 2/100
3375/3375          11s 2ms/step -
accuracy: 0.8985 - loss: 0.3770 - val_accuracy: 0.9352 - val_loss: 0.2546
Epoch 3/100
3375/3375          6s 2ms/step -
accuracy: 0.9168 - loss: 0.3005 - val_accuracy: 0.9407 - val_loss: 0.2211
Epoch 4/100
3375/3375          10s 2ms/step -
accuracy: 0.9248 - loss: 0.2652 - val_accuracy: 0.9445 - val_loss: 0.2018
Epoch 5/100
3375/3375          6s 2ms/step -
accuracy: 0.9330 - loss: 0.2371 - val_accuracy: 0.9482 - val_loss: 0.1890
Epoch 6/100
3375/3375          6s 2ms/step -
accuracy: 0.9377 - loss: 0.2219 - val_accuracy: 0.9523 - val_loss: 0.1766
Epoch 7/100
3375/3375          6s 2ms/step -
accuracy: 0.9409 - loss: 0.2069 - val_accuracy: 0.9545 - val_loss: 0.1661
Epoch 8/100
3375/3375          6s 2ms/step -
accuracy: 0.9453 - loss: 0.1932 - val_accuracy: 0.9565 - val_loss: 0.1612
Epoch 9/100
3375/3375          6s 2ms/step -
accuracy: 0.9461 - loss: 0.1894 - val_accuracy: 0.9570 - val_loss: 0.1540
Epoch 10/100
3375/3375          10s 2ms/step -
accuracy: 0.9489 - loss: 0.1757 - val_accuracy: 0.9590 - val_loss: 0.1467
Epoch 11/100
3375/3375          5s 2ms/step -
accuracy: 0.9529 - loss: 0.1641 - val_accuracy: 0.9610 - val_loss: 0.1443
Epoch 12/100

```

3375/3375                    6s 2ms/step -  
 accuracy: 0.9533 - loss: 0.1615 - val\_accuracy: 0.9613 - val\_loss: 0.1387  
 Epoch 13/100  
 3375/3375                    6s 2ms/step -  
 accuracy: 0.9560 - loss: 0.1531 - val\_accuracy: 0.9610 - val\_loss: 0.1373  
 Epoch 14/100  
 3375/3375                    6s 2ms/step -  
 accuracy: 0.9578 - loss: 0.1493 - val\_accuracy: 0.9630 - val\_loss: 0.1314  
 Epoch 15/100  
 3375/3375                    10s 2ms/step -  
 accuracy: 0.9611 - loss: 0.1377 - val\_accuracy: 0.9638 - val\_loss: 0.1289  
 Epoch 16/100  
 3375/3375                    10s 2ms/step -  
 accuracy: 0.9615 - loss: 0.1328 - val\_accuracy: 0.9633 - val\_loss: 0.1304  
 Epoch 17/100  
 3375/3375                    10s 2ms/step -  
 accuracy: 0.9633 - loss: 0.1255 - val\_accuracy: 0.9640 - val\_loss: 0.1271  
 Epoch 18/100  
 3375/3375                    6s 2ms/step -  
 accuracy: 0.9632 - loss: 0.1294 - val\_accuracy: 0.9658 - val\_loss: 0.1241  
 Epoch 19/100  
 3375/3375                    6s 2ms/step -  
 accuracy: 0.9646 - loss: 0.1231 - val\_accuracy: 0.9653 - val\_loss: 0.1235  
 Epoch 20/100  
 3375/3375                    10s 2ms/step -  
 accuracy: 0.9668 - loss: 0.1154 - val\_accuracy: 0.9665 - val\_loss: 0.1217  
 Epoch 21/100  
 3375/3375                    10s 2ms/step -  
 accuracy: 0.9684 - loss: 0.1089 - val\_accuracy: 0.9662 - val\_loss: 0.1233  
 Epoch 22/100  
 3375/3375                    6s 2ms/step -  
 accuracy: 0.9676 - loss: 0.1081 - val\_accuracy: 0.9662 - val\_loss: 0.1197  
 Epoch 23/100  
 3375/3375                    10s 2ms/step -  
 accuracy: 0.9687 - loss: 0.1092 - val\_accuracy: 0.9670 - val\_loss: 0.1202  
 Epoch 24/100  
 3375/3375                    10s 2ms/step -  
 accuracy: 0.9699 - loss: 0.1039 - val\_accuracy: 0.9670 - val\_loss: 0.1179  
 Epoch 25/100  
 3375/3375                    10s 2ms/step -  
 accuracy: 0.9717 - loss: 0.1001 - val\_accuracy: 0.9675 - val\_loss: 0.1175  
 Epoch 26/100  
 3375/3375                    10s 2ms/step -  
 accuracy: 0.9730 - loss: 0.0940 - val\_accuracy: 0.9678 - val\_loss: 0.1164  
 Epoch 27/100  
 3375/3375                    6s 2ms/step -  
 accuracy: 0.9714 - loss: 0.0980 - val\_accuracy: 0.9678 - val\_loss: 0.1141  
 Epoch 28/100

```

3375/3375          6s 2ms/step -
accuracy: 0.9723 - loss: 0.0940 - val_accuracy: 0.9683 - val_loss: 0.1161
Epoch 29/100
3375/3375          10s 2ms/step -
accuracy: 0.9719 - loss: 0.0921 - val_accuracy: 0.9690 - val_loss: 0.1146
Epoch 30/100
3375/3375          6s 2ms/step -
accuracy: 0.9744 - loss: 0.0875 - val_accuracy: 0.9678 - val_loss: 0.1178
Epoch 30: early stopping
Restoring model weights from the end of the best epoch: 27.
Training with batch size 16 and learning rate 1e-05
Epoch 1/100
3375/3375          6s 2ms/step -
accuracy: 0.2245 - loss: 2.1659 - val_accuracy: 0.5757 - val_loss: 1.6212
Epoch 2/100
3375/3375          11s 2ms/step -
accuracy: 0.6095 - loss: 1.5118 - val_accuracy: 0.7347 - val_loss: 1.1425
Epoch 3/100
3375/3375          10s 2ms/step -
accuracy: 0.7328 - loss: 1.1096 - val_accuracy: 0.8160 - val_loss: 0.8452
Epoch 4/100
3375/3375          10s 2ms/step -
accuracy: 0.7972 - loss: 0.8527 - val_accuracy: 0.8535 - val_loss: 0.6679
Epoch 5/100
3375/3375          10s 2ms/step -
accuracy: 0.8317 - loss: 0.7000 - val_accuracy: 0.8748 - val_loss: 0.5601
Epoch 6/100
3375/3375          10s 2ms/step -
accuracy: 0.8511 - loss: 0.6081 - val_accuracy: 0.8893 - val_loss: 0.4903
Epoch 7/100
3375/3375          6s 2ms/step -
accuracy: 0.8635 - loss: 0.5415 - val_accuracy: 0.8980 - val_loss: 0.4427
Epoch 8/100
3375/3375          6s 2ms/step -
accuracy: 0.8714 - loss: 0.5019 - val_accuracy: 0.9037 - val_loss: 0.4084
Epoch 9/100
3375/3375          10s 2ms/step -
accuracy: 0.8788 - loss: 0.4670 - val_accuracy: 0.9073 - val_loss: 0.3824
Epoch 10/100
3375/3375          6s 2ms/step -
accuracy: 0.8805 - loss: 0.4465 - val_accuracy: 0.9113 - val_loss: 0.3620
Epoch 11/100
3375/3375          10s 2ms/step -
accuracy: 0.8883 - loss: 0.4181 - val_accuracy: 0.9142 - val_loss: 0.3457
Epoch 12/100
3375/3375          10s 2ms/step -
accuracy: 0.8910 - loss: 0.4051 - val_accuracy: 0.9165 - val_loss: 0.3322
Epoch 13/100

```



3375/3375                    6s 2ms/step -  
 accuracy: 0.8930 - loss: 0.3923 - val\_accuracy: 0.9175 - val\_loss: 0.3209  
 Epoch 14/100  
 3375/3375                    6s 2ms/step -  
 accuracy: 0.8935 - loss: 0.3835 - val\_accuracy: 0.9192 - val\_loss: 0.3114  
 Epoch 15/100  
 3375/3375                    6s 2ms/step -  
 accuracy: 0.9001 - loss: 0.3677 - val\_accuracy: 0.9207 - val\_loss: 0.3031  
 Epoch 16/100  
 3375/3375                    10s 2ms/step -  
 accuracy: 0.9017 - loss: 0.3568 - val\_accuracy: 0.9205 - val\_loss: 0.2955  
 Epoch 17/100  
 3375/3375                    6s 2ms/step -  
 accuracy: 0.9037 - loss: 0.3500 - val\_accuracy: 0.9222 - val\_loss: 0.2889  
 Epoch 18/100  
 3375/3375                    6s 2ms/step -  
 accuracy: 0.9050 - loss: 0.3417 - val\_accuracy: 0.9238 - val\_loss: 0.2833  
 Epoch 19/100  
 3375/3375                    6s 2ms/step -  
 accuracy: 0.9088 - loss: 0.3298 - val\_accuracy: 0.9248 - val\_loss: 0.2781  
 Epoch 20/100  
 3375/3375                    6s 2ms/step -  
 accuracy: 0.9059 - loss: 0.3340 - val\_accuracy: 0.9258 - val\_loss: 0.2727  
 Epoch 21/100  
 3375/3375                    6s 2ms/step -  
 accuracy: 0.9092 - loss: 0.3242 - val\_accuracy: 0.9265 - val\_loss: 0.2684  
 Epoch 22/100  
 3375/3375                    10s 2ms/step -  
 accuracy: 0.9098 - loss: 0.3206 - val\_accuracy: 0.9290 - val\_loss: 0.2642  
 Epoch 23/100  
 3375/3375                    10s 2ms/step -  
 accuracy: 0.9104 - loss: 0.3171 - val\_accuracy: 0.9292 - val\_loss: 0.2607  
 Epoch 24/100  
 3375/3375                    10s 2ms/step -  
 accuracy: 0.9118 - loss: 0.3108 - val\_accuracy: 0.9302 - val\_loss: 0.2570  
 Epoch 25/100  
 3375/3375                    11s 2ms/step -  
 accuracy: 0.9145 - loss: 0.3069 - val\_accuracy: 0.9297 - val\_loss: 0.2540  
 Epoch 26/100  
 3375/3375                    5s 2ms/step -  
 accuracy: 0.9141 - loss: 0.3030 - val\_accuracy: 0.9300 - val\_loss: 0.2509  
 Epoch 27/100  
 3375/3375                    6s 2ms/step -  
 accuracy: 0.9145 - loss: 0.3000 - val\_accuracy: 0.9312 - val\_loss: 0.2476  
 Epoch 28/100  
 3375/3375                    6s 2ms/step -  
 accuracy: 0.9153 - loss: 0.2969 - val\_accuracy: 0.9310 - val\_loss: 0.2453  
 Epoch 29/100

3375/3375                    6s 2ms/step -  
accuracy: 0.9171 - loss: 0.2917 - val\_accuracy: 0.9333 - val\_loss: 0.2426  
Epoch 30/100

3375/3375                    6s 2ms/step -  
accuracy: 0.9177 - loss: 0.2892 - val\_accuracy: 0.9333 - val\_loss: 0.2401  
Epoch 31/100

3375/3375                    10s 2ms/step -  
accuracy: 0.9187 - loss: 0.2937 - val\_accuracy: 0.9338 - val\_loss: 0.2375  
Epoch 32/100

3375/3375                    6s 2ms/step -  
accuracy: 0.9196 - loss: 0.2842 - val\_accuracy: 0.9340 - val\_loss: 0.2358  
Epoch 33/100

3375/3375                    6s 2ms/step -  
accuracy: 0.9195 - loss: 0.2821 - val\_accuracy: 0.9350 - val\_loss: 0.2333  
Epoch 34/100

3375/3375                    6s 2ms/step -  
accuracy: 0.9210 - loss: 0.2758 - val\_accuracy: 0.9350 - val\_loss: 0.2312  
Epoch 35/100

3375/3375                    6s 2ms/step -  
accuracy: 0.9217 - loss: 0.2777 - val\_accuracy: 0.9357 - val\_loss: 0.2295  
Epoch 36/100

3375/3375                    10s 2ms/step -  
accuracy: 0.9222 - loss: 0.2755 - val\_accuracy: 0.9382 - val\_loss: 0.2277  
Epoch 37/100

3375/3375                    6s 2ms/step -  
accuracy: 0.9234 - loss: 0.2687 - val\_accuracy: 0.9382 - val\_loss: 0.2257  
Epoch 38/100

3375/3375                    10s 2ms/step -  
accuracy: 0.9239 - loss: 0.2709 - val\_accuracy: 0.9377 - val\_loss: 0.2240  
Epoch 39/100

3375/3375                    5s 2ms/step -  
accuracy: 0.9252 - loss: 0.2664 - val\_accuracy: 0.9365 - val\_loss: 0.2230  
Epoch 40/100

3375/3375                    11s 2ms/step -  
accuracy: 0.9239 - loss: 0.2642 - val\_accuracy: 0.9385 - val\_loss: 0.2206  
Epoch 41/100

3375/3375                    10s 2ms/step -  
accuracy: 0.9252 - loss: 0.2606 - val\_accuracy: 0.9378 - val\_loss: 0.2194  
Epoch 42/100

3375/3375                    6s 2ms/step -  
accuracy: 0.9273 - loss: 0.2538 - val\_accuracy: 0.9392 - val\_loss: 0.2174  
Epoch 43/100

3375/3375                    10s 2ms/step -  
accuracy: 0.9268 - loss: 0.2596 - val\_accuracy: 0.9385 - val\_loss: 0.2165  
Epoch 44/100

3375/3375                    10s 2ms/step -  
accuracy: 0.9277 - loss: 0.2505 - val\_accuracy: 0.9395 - val\_loss: 0.2149  
Epoch 45/100

3375/3375                    11s 2ms/step -  
 accuracy: 0.9266 - loss: 0.2586 - val\_accuracy: 0.9392 - val\_loss: 0.2133  
 Epoch 46/100  
 3375/3375                    6s 2ms/step -  
 accuracy: 0.9267 - loss: 0.2541 - val\_accuracy: 0.9400 - val\_loss: 0.2123  
 Epoch 47/100  
 3375/3375                    10s 2ms/step -  
 accuracy: 0.9267 - loss: 0.2596 - val\_accuracy: 0.9403 - val\_loss: 0.2103  
 Epoch 48/100  
 3375/3375                    6s 2ms/step -  
 accuracy: 0.9265 - loss: 0.2528 - val\_accuracy: 0.9405 - val\_loss: 0.2094  
 Epoch 49/100  
 3375/3375                    10s 2ms/step -  
 accuracy: 0.9278 - loss: 0.2530 - val\_accuracy: 0.9412 - val\_loss: 0.2077  
 Epoch 50/100  
 3375/3375                    10s 2ms/step -  
 accuracy: 0.9308 - loss: 0.2431 - val\_accuracy: 0.9413 - val\_loss: 0.2065  
 Epoch 51/100  
 3375/3375                    6s 2ms/step -  
 accuracy: 0.9314 - loss: 0.2454 - val\_accuracy: 0.9420 - val\_loss: 0.2055  
 Epoch 52/100  
 3375/3375                    6s 2ms/step -  
 accuracy: 0.9304 - loss: 0.2458 - val\_accuracy: 0.9420 - val\_loss: 0.2046  
 Epoch 53/100  
 3375/3375                    10s 2ms/step -  
 accuracy: 0.9310 - loss: 0.2450 - val\_accuracy: 0.9425 - val\_loss: 0.2028  
 Epoch 54/100  
 3375/3375                    11s 2ms/step -  
 accuracy: 0.9313 - loss: 0.2452 - val\_accuracy: 0.9427 - val\_loss: 0.2018  
 Epoch 55/100  
 3375/3375                    10s 2ms/step -  
 accuracy: 0.9322 - loss: 0.2398 - val\_accuracy: 0.9435 - val\_loss: 0.2007  
 Epoch 56/100  
 3375/3375                    6s 2ms/step -  
 accuracy: 0.9315 - loss: 0.2359 - val\_accuracy: 0.9448 - val\_loss: 0.1993  
 Epoch 57/100  
 3375/3375                    10s 2ms/step -  
 accuracy: 0.9349 - loss: 0.2355 - val\_accuracy: 0.9453 - val\_loss: 0.1981  
 Epoch 58/100  
 3375/3375                    6s 2ms/step -  
 accuracy: 0.9337 - loss: 0.2337 - val\_accuracy: 0.9457 - val\_loss: 0.1972  
 Epoch 59/100  
 3375/3375                    10s 2ms/step -  
 accuracy: 0.9331 - loss: 0.2327 - val\_accuracy: 0.9452 - val\_loss: 0.1960  
 Epoch 60/100  
 3375/3375                    6s 2ms/step -  
 accuracy: 0.9349 - loss: 0.2303 - val\_accuracy: 0.9458 - val\_loss: 0.1952  
 Epoch 61/100

3375/3375                    10s 2ms/step -  
accuracy: 0.9350 - loss: 0.2290 - val\_accuracy: 0.9462 - val\_loss: 0.1941  
Epoch 62/100

3375/3375                    11s 2ms/step -  
accuracy: 0.9365 - loss: 0.2210 - val\_accuracy: 0.9468 - val\_loss: 0.1929  
Epoch 63/100

3375/3375                    10s 2ms/step -  
accuracy: 0.9342 - loss: 0.2288 - val\_accuracy: 0.9480 - val\_loss: 0.1916  
Epoch 64/100

3375/3375                    10s 2ms/step -  
accuracy: 0.9354 - loss: 0.2258 - val\_accuracy: 0.9472 - val\_loss: 0.1904  
Epoch 65/100

3375/3375                    10s 2ms/step -  
accuracy: 0.9369 - loss: 0.2200 - val\_accuracy: 0.9480 - val\_loss: 0.1901  
Epoch 66/100

3375/3375                    10s 2ms/step -  
accuracy: 0.9383 - loss: 0.2156 - val\_accuracy: 0.9488 - val\_loss: 0.1886  
Epoch 67/100

3375/3375                    10s 2ms/step -  
accuracy: 0.9381 - loss: 0.2176 - val\_accuracy: 0.9475 - val\_loss: 0.1879  
Epoch 68/100

3375/3375                    10s 2ms/step -  
accuracy: 0.9352 - loss: 0.2255 - val\_accuracy: 0.9490 - val\_loss: 0.1865  
Epoch 69/100

3375/3375                    6s 2ms/step -  
accuracy: 0.9372 - loss: 0.2173 - val\_accuracy: 0.9485 - val\_loss: 0.1859  
Epoch 70/100

3375/3375                    6s 2ms/step -  
accuracy: 0.9379 - loss: 0.2160 - val\_accuracy: 0.9485 - val\_loss: 0.1850  
Epoch 71/100

3375/3375                    6s 2ms/step -  
accuracy: 0.9357 - loss: 0.2186 - val\_accuracy: 0.9483 - val\_loss: 0.1841  
Epoch 72/100

3375/3375                    5s 2ms/step -  
accuracy: 0.9406 - loss: 0.2072 - val\_accuracy: 0.9482 - val\_loss: 0.1837  
Epoch 73/100

3375/3375                    11s 2ms/step -  
accuracy: 0.9395 - loss: 0.2100 - val\_accuracy: 0.9488 - val\_loss: 0.1824  
Epoch 74/100

3375/3375                    6s 2ms/step -  
accuracy: 0.9376 - loss: 0.2134 - val\_accuracy: 0.9498 - val\_loss: 0.1814  
Epoch 75/100

3375/3375                    10s 2ms/step -  
accuracy: 0.9410 - loss: 0.2050 - val\_accuracy: 0.9493 - val\_loss: 0.1804  
Epoch 76/100

3375/3375                    6s 2ms/step -  
accuracy: 0.9418 - loss: 0.2040 - val\_accuracy: 0.9497 - val\_loss: 0.1796  
Epoch 77/100

3375/3375                    6s 2ms/step -  
accuracy: 0.9406 - loss: 0.2064 - val\_accuracy: 0.9490 - val\_loss: 0.1795  
Epoch 78/100

3375/3375                    6s 2ms/step -  
accuracy: 0.9418 - loss: 0.2050 - val\_accuracy: 0.9490 - val\_loss: 0.1780  
Epoch 79/100

3375/3375                    10s 2ms/step -  
accuracy: 0.9412 - loss: 0.2032 - val\_accuracy: 0.9497 - val\_loss: 0.1772  
Epoch 80/100

3375/3375                    10s 2ms/step -  
accuracy: 0.9410 - loss: 0.2052 - val\_accuracy: 0.9500 - val\_loss: 0.1763  
Epoch 81/100

3375/3375                    6s 2ms/step -  
accuracy: 0.9429 - loss: 0.2002 - val\_accuracy: 0.9505 - val\_loss: 0.1759  
Epoch 82/100

3375/3375                    10s 2ms/step -  
accuracy: 0.9425 - loss: 0.2023 - val\_accuracy: 0.9500 - val\_loss: 0.1751  
Epoch 83/100

3375/3375                    6s 2ms/step -  
accuracy: 0.9422 - loss: 0.1977 - val\_accuracy: 0.9502 - val\_loss: 0.1740  
Epoch 84/100

3375/3375                    6s 2ms/step -  
accuracy: 0.9435 - loss: 0.1961 - val\_accuracy: 0.9513 - val\_loss: 0.1734  
Epoch 85/100

3375/3375                    6s 2ms/step -  
accuracy: 0.9428 - loss: 0.1994 - val\_accuracy: 0.9510 - val\_loss: 0.1729  
Epoch 86/100

3375/3375                    6s 2ms/step -  
accuracy: 0.9437 - loss: 0.1949 - val\_accuracy: 0.9513 - val\_loss: 0.1720  
Epoch 87/100

3375/3375                    10s 2ms/step -  
accuracy: 0.9427 - loss: 0.1960 - val\_accuracy: 0.9513 - val\_loss: 0.1711  
Epoch 88/100

3375/3375                    10s 2ms/step -  
accuracy: 0.9443 - loss: 0.1969 - val\_accuracy: 0.9520 - val\_loss: 0.1706  
Epoch 89/100

3375/3375                    6s 2ms/step -  
accuracy: 0.9445 - loss: 0.1926 - val\_accuracy: 0.9515 - val\_loss: 0.1695  
Epoch 90/100

3375/3375                    10s 2ms/step -  
accuracy: 0.9460 - loss: 0.1884 - val\_accuracy: 0.9522 - val\_loss: 0.1691  
Epoch 91/100

3375/3375                    10s 2ms/step -  
accuracy: 0.9455 - loss: 0.1897 - val\_accuracy: 0.9518 - val\_loss: 0.1687  
Epoch 92/100

3375/3375                    10s 2ms/step -  
accuracy: 0.9447 - loss: 0.1943 - val\_accuracy: 0.9520 - val\_loss: 0.1683  
Epoch 93/100



```

3375/3375          6s 2ms/step -
accuracy: 0.9441 - loss: 0.1941 - val_accuracy: 0.9528 - val_loss: 0.1675
Epoch 94/100
3375/3375          6s 2ms/step -
accuracy: 0.9451 - loss: 0.1901 - val_accuracy: 0.9527 - val_loss: 0.1663
Epoch 95/100
3375/3375          10s 2ms/step -
accuracy: 0.9461 - loss: 0.1876 - val_accuracy: 0.9530 - val_loss: 0.1658
Epoch 96/100
3375/3375          6s 2ms/step -
accuracy: 0.9450 - loss: 0.1904 - val_accuracy: 0.9535 - val_loss: 0.1652
Epoch 97/100
3375/3375          10s 2ms/step -
accuracy: 0.9463 - loss: 0.1879 - val_accuracy: 0.9527 - val_loss: 0.1646
Epoch 98/100
3375/3375          6s 2ms/step -
accuracy: 0.9483 - loss: 0.1843 - val_accuracy: 0.9523 - val_loss: 0.1641
Epoch 99/100
3375/3375          6s 2ms/step -
accuracy: 0.9469 - loss: 0.1865 - val_accuracy: 0.9527 - val_loss: 0.1639
Epoch 100/100
3375/3375          10s 2ms/step -
accuracy: 0.9467 - loss: 0.1844 - val_accuracy: 0.9537 - val_loss: 0.1631
Restoring model weights from the end of the best epoch: 100.
Training with batch size 32 and learning rate 0.01
Epoch 1/100
1688/1688          4s 2ms/step -
accuracy: 0.8554 - loss: 0.4671 - val_accuracy: 0.9340 - val_loss: 0.2196
Epoch 2/100
1688/1688          5s 2ms/step -
accuracy: 0.9379 - loss: 0.2108 - val_accuracy: 0.9503 - val_loss: 0.1874
Epoch 3/100
1688/1688          3s 2ms/step -
accuracy: 0.9465 - loss: 0.1868 - val_accuracy: 0.9575 - val_loss: 0.1558
Epoch 4/100
1688/1688          4s 2ms/step -
accuracy: 0.9477 - loss: 0.1812 - val_accuracy: 0.9553 - val_loss: 0.1687
Epoch 5/100
1688/1688          3s 2ms/step -
accuracy: 0.9522 - loss: 0.1627 - val_accuracy: 0.9602 - val_loss: 0.1454
Epoch 6/100
1688/1688          3s 2ms/step -
accuracy: 0.9564 - loss: 0.1545 - val_accuracy: 0.9573 - val_loss: 0.1534
Epoch 7/100
1688/1688          5s 2ms/step -
accuracy: 0.9581 - loss: 0.1487 - val_accuracy: 0.9590 - val_loss: 0.1443
Epoch 8/100
1688/1688          3s 2ms/step -

```

```

accuracy: 0.9592 - loss: 0.1462 - val_accuracy: 0.9587 - val_loss: 0.1570
Epoch 9/100
1688/1688          5s 2ms/step -
accuracy: 0.9608 - loss: 0.1402 - val_accuracy: 0.9582 - val_loss: 0.1574
Epoch 10/100
1688/1688          4s 2ms/step -
accuracy: 0.9627 - loss: 0.1354 - val_accuracy: 0.9667 - val_loss: 0.1328
Epoch 11/100
1688/1688          5s 2ms/step -
accuracy: 0.9623 - loss: 0.1331 - val_accuracy: 0.9640 - val_loss: 0.1452
Epoch 12/100
1688/1688          3s 2ms/step -
accuracy: 0.9623 - loss: 0.1326 - val_accuracy: 0.9590 - val_loss: 0.1815
Epoch 13/100
1688/1688          6s 2ms/step -
accuracy: 0.9619 - loss: 0.1385 - val_accuracy: 0.9535 - val_loss: 0.1928
Epoch 13: early stopping
Restoring model weights from the end of the best epoch: 10.
Training with batch size 32 and learning rate 0.001
Epoch 1/100
1688/1688          4s 2ms/step -
accuracy: 0.7850 - loss: 0.7054 - val_accuracy: 0.9493 - val_loss: 0.1786
Epoch 2/100
1688/1688          6s 2ms/step -
accuracy: 0.9422 - loss: 0.1968 - val_accuracy: 0.9597 - val_loss: 0.1457
Epoch 3/100
1688/1688          4s 2ms/step -
accuracy: 0.9551 - loss: 0.1534 - val_accuracy: 0.9638 - val_loss: 0.1319
Epoch 4/100
1688/1688          5s 2ms/step -
accuracy: 0.9618 - loss: 0.1287 - val_accuracy: 0.9688 - val_loss: 0.1141
Epoch 5/100
1688/1688          4s 2ms/step -
accuracy: 0.9668 - loss: 0.1095 - val_accuracy: 0.9643 - val_loss: 0.1144
Epoch 6/100
1688/1688          3s 2ms/step -
accuracy: 0.9716 - loss: 0.0893 - val_accuracy: 0.9658 - val_loss: 0.1151
Epoch 7/100
1688/1688          3s 2ms/step -
accuracy: 0.9731 - loss: 0.0860 - val_accuracy: 0.9693 - val_loss: 0.1082
Epoch 8/100
1688/1688          5s 2ms/step -
accuracy: 0.9757 - loss: 0.0780 - val_accuracy: 0.9683 - val_loss: 0.1093
Epoch 9/100
1688/1688          3s 2ms/step -
accuracy: 0.9792 - loss: 0.0671 - val_accuracy: 0.9647 - val_loss: 0.1225
Epoch 10/100
1688/1688          3s 2ms/step -

```

accuracy: 0.9805 - loss: 0.0649 - val\_accuracy: 0.9683 - val\_loss: 0.1134  
 Epoch 10: early stopping  
 Restoring model weights from the end of the best epoch: 7.  
 Training with batch size 32 and learning rate 0.0001  
 Epoch 1/100  
 1688/1688                    4s 2ms/step -  
 accuracy: 0.4932 - loss: 1.6697 - val\_accuracy: 0.8653 - val\_loss: 0.5202  
 Epoch 2/100  
 1688/1688                    5s 2ms/step -  
 accuracy: 0.8603 - loss: 0.5166 - val\_accuracy: 0.9180 - val\_loss: 0.3119  
 Epoch 3/100  
 1688/1688                    4s 2ms/step -  
 accuracy: 0.8995 - loss: 0.3582 - val\_accuracy: 0.9322 - val\_loss: 0.2480  
 Epoch 4/100  
 1688/1688                    5s 2ms/step -  
 accuracy: 0.9145 - loss: 0.2948 - val\_accuracy: 0.9397 - val\_loss: 0.2187  
 Epoch 5/100  
 1688/1688                    3s 2ms/step -  
 accuracy: 0.9251 - loss: 0.2605 - val\_accuracy: 0.9443 - val\_loss: 0.1997  
 Epoch 6/100  
 1688/1688                    4s 2ms/step -  
 accuracy: 0.9320 - loss: 0.2368 - val\_accuracy: 0.9490 - val\_loss: 0.1876  
 Epoch 7/100  
 1688/1688                    3s 2ms/step -  
 accuracy: 0.9383 - loss: 0.2135 - val\_accuracy: 0.9538 - val\_loss: 0.1763  
 Epoch 8/100  
 1688/1688                    3s 2ms/step -  
 accuracy: 0.9395 - loss: 0.2100 - val\_accuracy: 0.9550 - val\_loss: 0.1673  
 Epoch 9/100  
 1688/1688                    4s 2ms/step -  
 accuracy: 0.9436 - loss: 0.1980 - val\_accuracy: 0.9573 - val\_loss: 0.1619  
 Epoch 10/100  
 1688/1688                    3s 2ms/step -  
 accuracy: 0.9453 - loss: 0.1891 - val\_accuracy: 0.9582 - val\_loss: 0.1555  
 Epoch 11/100  
 1688/1688                    5s 2ms/step -  
 accuracy: 0.9492 - loss: 0.1771 - val\_accuracy: 0.9598 - val\_loss: 0.1520  
 Epoch 12/100  
 1688/1688                    4s 2ms/step -  
 accuracy: 0.9519 - loss: 0.1706 - val\_accuracy: 0.9615 - val\_loss: 0.1446  
 Epoch 13/100  
 1688/1688                    5s 2ms/step -  
 accuracy: 0.9538 - loss: 0.1585 - val\_accuracy: 0.9637 - val\_loss: 0.1415  
 Epoch 14/100  
 1688/1688                    5s 2ms/step -  
 accuracy: 0.9557 - loss: 0.1558 - val\_accuracy: 0.9638 - val\_loss: 0.1385  
 Epoch 15/100  
 1688/1688                    5s 2ms/step -

accuracy: 0.9559 - loss: 0.1513 - val\_accuracy: 0.9633 - val\_loss: 0.1352  
 Epoch 16/100  
 1688/1688                    3s 2ms/step -  
 accuracy: 0.9573 - loss: 0.1449 - val\_accuracy: 0.9652 - val\_loss: 0.1314  
 Epoch 17/100  
 1688/1688                    5s 2ms/step -  
 accuracy: 0.9601 - loss: 0.1401 - val\_accuracy: 0.9673 - val\_loss: 0.1302  
 Epoch 18/100  
 1688/1688                    5s 2ms/step -  
 accuracy: 0.9614 - loss: 0.1354 - val\_accuracy: 0.9673 - val\_loss: 0.1283  
 Epoch 19/100  
 1688/1688                    6s 2ms/step -  
 accuracy: 0.9609 - loss: 0.1321 - val\_accuracy: 0.9673 - val\_loss: 0.1257  
 Epoch 20/100  
 1688/1688                    3s 2ms/step -  
 accuracy: 0.9621 - loss: 0.1325 - val\_accuracy: 0.9675 - val\_loss: 0.1238  
 Epoch 21/100  
 1688/1688                    5s 2ms/step -  
 accuracy: 0.9641 - loss: 0.1231 - val\_accuracy: 0.9677 - val\_loss: 0.1232  
 Epoch 22/100  
 1688/1688                    6s 2ms/step -  
 accuracy: 0.9628 - loss: 0.1255 - val\_accuracy: 0.9677 - val\_loss: 0.1204  
 Epoch 23/100  
 1688/1688                    3s 2ms/step -  
 accuracy: 0.9646 - loss: 0.1203 - val\_accuracy: 0.9670 - val\_loss: 0.1210  
 Epoch 24/100  
 1688/1688                    3s 2ms/step -  
 accuracy: 0.9660 - loss: 0.1139 - val\_accuracy: 0.9670 - val\_loss: 0.1188  
 Epoch 25/100  
 1688/1688                    4s 2ms/step -  
 accuracy: 0.9669 - loss: 0.1142 - val\_accuracy: 0.9680 - val\_loss: 0.1170  
 Epoch 26/100  
 1688/1688                    4s 2ms/step -  
 accuracy: 0.9697 - loss: 0.1065 - val\_accuracy: 0.9678 - val\_loss: 0.1150  
 Epoch 27/100  
 1688/1688                    6s 2ms/step -  
 accuracy: 0.9697 - loss: 0.1059 - val\_accuracy: 0.9677 - val\_loss: 0.1153  
 Epoch 28/100  
 1688/1688                    5s 2ms/step -  
 accuracy: 0.9684 - loss: 0.1044 - val\_accuracy: 0.9670 - val\_loss: 0.1161  
 Epoch 29/100  
 1688/1688                    3s 2ms/step -  
 accuracy: 0.9701 - loss: 0.1022 - val\_accuracy: 0.9675 - val\_loss: 0.1142  
 Epoch 30/100  
 1688/1688                    4s 2ms/step -  
 accuracy: 0.9713 - loss: 0.0976 - val\_accuracy: 0.9677 - val\_loss: 0.1131  
 Epoch 31/100  
 1688/1688                    3s 2ms/step -

accuracy: 0.9716 - loss: 0.0979 - val\_accuracy: 0.9680 - val\_loss: 0.1125  
 Epoch 32/100  
 1688/1688 5s 2ms/step -  
 accuracy: 0.9724 - loss: 0.0948 - val\_accuracy: 0.9678 - val\_loss: 0.1116  
 Epoch 33/100  
 1688/1688 6s 2ms/step -  
 accuracy: 0.9725 - loss: 0.0936 - val\_accuracy: 0.9680 - val\_loss: 0.1122  
 Epoch 34/100  
 1688/1688 4s 2ms/step -  
 accuracy: 0.9728 - loss: 0.0906 - val\_accuracy: 0.9690 - val\_loss: 0.1106  
 Epoch 35/100  
 1688/1688 6s 2ms/step -  
 accuracy: 0.9752 - loss: 0.0870 - val\_accuracy: 0.9672 - val\_loss: 0.1102  
 Epoch 36/100  
 1688/1688 3s 2ms/step -  
 accuracy: 0.9741 - loss: 0.0897 - val\_accuracy: 0.9693 - val\_loss: 0.1100  
 Epoch 37/100  
 1688/1688 3s 2ms/step -  
 accuracy: 0.9727 - loss: 0.0899 - val\_accuracy: 0.9687 - val\_loss: 0.1107  
 Epoch 38/100  
 1688/1688 3s 2ms/step -  
 accuracy: 0.9755 - loss: 0.0849 - val\_accuracy: 0.9693 - val\_loss: 0.1082  
 Epoch 39/100  
 1688/1688 4s 2ms/step -  
 accuracy: 0.9776 - loss: 0.0821 - val\_accuracy: 0.9705 - val\_loss: 0.1085  
 Epoch 40/100  
 1688/1688 4s 2ms/step -  
 accuracy: 0.9761 - loss: 0.0830 - val\_accuracy: 0.9702 - val\_loss: 0.1076  
 Epoch 41/100  
 1688/1688 6s 2ms/step -  
 accuracy: 0.9773 - loss: 0.0781 - val\_accuracy: 0.9698 - val\_loss: 0.1076  
 Epoch 42/100  
 1688/1688 3s 2ms/step -  
 accuracy: 0.9774 - loss: 0.0792 - val\_accuracy: 0.9690 - val\_loss: 0.1087  
 Epoch 43/100  
 1688/1688 3s 2ms/step -  
 accuracy: 0.9769 - loss: 0.0813 - val\_accuracy: 0.9710 - val\_loss: 0.1068  
 Epoch 44/100  
 1688/1688 3s 2ms/step -  
 accuracy: 0.9783 - loss: 0.0774 - val\_accuracy: 0.9707 - val\_loss: 0.1070  
 Epoch 45/100  
 1688/1688 5s 2ms/step -  
 accuracy: 0.9781 - loss: 0.0762 - val\_accuracy: 0.9702 - val\_loss: 0.1081  
 Epoch 46/100  
 1688/1688 5s 2ms/step -  
 accuracy: 0.9792 - loss: 0.0750 - val\_accuracy: 0.9697 - val\_loss: 0.1076  
 Epoch 46: early stopping  
 Restoring model weights from the end of the best epoch: 43.

Training with batch size 32 and learning rate 1e-05

Epoch 1/100

1688/1688 4s 2ms/step -

accuracy: 0.1859 - loss: 2.2611 - val\_accuracy: 0.4627 - val\_loss: 1.9556

Epoch 2/100

1688/1688 5s 2ms/step -

accuracy: 0.5143 - loss: 1.8525 - val\_accuracy: 0.6697 - val\_loss: 1.5229

Epoch 3/100

1688/1688 6s 2ms/step -

accuracy: 0.6635 - loss: 1.4635 - val\_accuracy: 0.7660 - val\_loss: 1.1992

Epoch 4/100

1688/1688 4s 2ms/step -

accuracy: 0.7390 - loss: 1.1777 - val\_accuracy: 0.8185 - val\_loss: 0.9495

Epoch 5/100

1688/1688 3s 2ms/step -

accuracy: 0.7882 - loss: 0.9635 - val\_accuracy: 0.8507 - val\_loss: 0.7750

Epoch 6/100

1688/1688 4s 2ms/step -

accuracy: 0.8151 - loss: 0.8107 - val\_accuracy: 0.8718 - val\_loss: 0.6566

Epoch 7/100

1688/1688 5s 2ms/step -

accuracy: 0.8385 - loss: 0.6997 - val\_accuracy: 0.8858 - val\_loss: 0.5730

Epoch 8/100

1688/1688 5s 2ms/step -

accuracy: 0.8525 - loss: 0.6227 - val\_accuracy: 0.8935 - val\_loss: 0.5128

Epoch 9/100

1688/1688 5s 2ms/step -

accuracy: 0.8577 - loss: 0.5780 - val\_accuracy: 0.9003 - val\_loss: 0.4674

Epoch 10/100

1688/1688 3s 2ms/step -

accuracy: 0.8697 - loss: 0.5311 - val\_accuracy: 0.9058 - val\_loss: 0.4322

Epoch 11/100

1688/1688 6s 2ms/step -

accuracy: 0.8773 - loss: 0.4944 - val\_accuracy: 0.9087 - val\_loss: 0.4041

Epoch 12/100

1688/1688 3s 2ms/step -

accuracy: 0.8841 - loss: 0.4665 - val\_accuracy: 0.9138 - val\_loss: 0.3814

Epoch 13/100

1688/1688 5s 2ms/step -

accuracy: 0.8870 - loss: 0.4439 - val\_accuracy: 0.9165 - val\_loss: 0.3623

Epoch 14/100

1688/1688 4s 2ms/step -

accuracy: 0.8918 - loss: 0.4260 - val\_accuracy: 0.9197 - val\_loss: 0.3462

Epoch 15/100

1688/1688 5s 2ms/step -

accuracy: 0.8965 - loss: 0.4062 - val\_accuracy: 0.9220 - val\_loss: 0.3323

Epoch 16/100

1688/1688 3s 2ms/step -



accuracy: 0.8976 - loss: 0.3919 - val\_accuracy: 0.9232 - val\_loss: 0.3205  
 Epoch 17/100  
 1688/1688                    4s 2ms/step -  
 accuracy: 0.8992 - loss: 0.3860 - val\_accuracy: 0.9245 - val\_loss: 0.3098  
 Epoch 18/100  
 1688/1688                    5s 2ms/step -  
 accuracy: 0.9030 - loss: 0.3747 - val\_accuracy: 0.9273 - val\_loss: 0.3005  
 Epoch 19/100  
 1688/1688                    3s 2ms/step -  
 accuracy: 0.9034 - loss: 0.3651 - val\_accuracy: 0.9280 - val\_loss: 0.2922  
 Epoch 20/100  
 1688/1688                    4s 2ms/step -  
 accuracy: 0.9079 - loss: 0.3518 - val\_accuracy: 0.9293 - val\_loss: 0.2845  
 Epoch 21/100  
 1688/1688                    3s 2ms/step -  
 accuracy: 0.9083 - loss: 0.3462 - val\_accuracy: 0.9288 - val\_loss: 0.2777  
 Epoch 22/100  
 1688/1688                    5s 2ms/step -  
 accuracy: 0.9095 - loss: 0.3381 - val\_accuracy: 0.9305 - val\_loss: 0.2715  
 Epoch 23/100  
 1688/1688                    5s 2ms/step -  
 accuracy: 0.9106 - loss: 0.3289 - val\_accuracy: 0.9310 - val\_loss: 0.2656  
 Epoch 24/100  
 1688/1688                    3s 2ms/step -  
 accuracy: 0.9104 - loss: 0.3289 - val\_accuracy: 0.9315 - val\_loss: 0.2606  
 Epoch 25/100  
 1688/1688                    6s 2ms/step -  
 accuracy: 0.9132 - loss: 0.3187 - val\_accuracy: 0.9335 - val\_loss: 0.2557  
 Epoch 26/100  
 1688/1688                    3s 2ms/step -  
 accuracy: 0.9151 - loss: 0.3134 - val\_accuracy: 0.9348 - val\_loss: 0.2513  
 Epoch 27/100  
 1688/1688                    3s 2ms/step -  
 accuracy: 0.9147 - loss: 0.3117 - val\_accuracy: 0.9355 - val\_loss: 0.2470  
 Epoch 28/100  
 1688/1688                    3s 2ms/step -  
 accuracy: 0.9179 - loss: 0.3038 - val\_accuracy: 0.9363 - val\_loss: 0.2430  
 Epoch 29/100  
 1688/1688                    5s 2ms/step -  
 accuracy: 0.9182 - loss: 0.3035 - val\_accuracy: 0.9363 - val\_loss: 0.2394  
 Epoch 30/100  
 1688/1688                    5s 2ms/step -  
 accuracy: 0.9198 - loss: 0.2922 - val\_accuracy: 0.9383 - val\_loss: 0.2359  
 Epoch 31/100  
 1688/1688                    4s 2ms/step -  
 accuracy: 0.9204 - loss: 0.2876 - val\_accuracy: 0.9390 - val\_loss: 0.2325  
 Epoch 32/100  
 1688/1688                    5s 2ms/step -

accuracy: 0.9221 - loss: 0.2825 - val\_accuracy: 0.9398 - val\_loss: 0.2292  
 Epoch 33/100  
 1688/1688                    3s 2ms/step -  
 accuracy: 0.9237 - loss: 0.2780 - val\_accuracy: 0.9403 - val\_loss: 0.2261  
 Epoch 34/100  
 1688/1688                    5s 2ms/step -  
 accuracy: 0.9245 - loss: 0.2761 - val\_accuracy: 0.9413 - val\_loss: 0.2234  
 Epoch 35/100  
 1688/1688                    3s 2ms/step -  
 accuracy: 0.9215 - loss: 0.2820 - val\_accuracy: 0.9422 - val\_loss: 0.2209  
 Epoch 36/100  
 1688/1688                    3s 2ms/step -  
 accuracy: 0.9259 - loss: 0.2709 - val\_accuracy: 0.9435 - val\_loss: 0.2182  
 Epoch 37/100  
 1688/1688                    4s 2ms/step -  
 accuracy: 0.9238 - loss: 0.2769 - val\_accuracy: 0.9438 - val\_loss: 0.2156  
 Epoch 38/100  
 1688/1688                    5s 2ms/step -  
 accuracy: 0.9267 - loss: 0.2676 - val\_accuracy: 0.9435 - val\_loss: 0.2132  
 Epoch 39/100  
 1688/1688                    6s 2ms/step -  
 accuracy: 0.9278 - loss: 0.2648 - val\_accuracy: 0.9443 - val\_loss: 0.2112  
 Epoch 40/100  
 1688/1688                    4s 2ms/step -  
 accuracy: 0.9285 - loss: 0.2594 - val\_accuracy: 0.9453 - val\_loss: 0.2087  
 Epoch 41/100  
 1688/1688                    5s 2ms/step -  
 accuracy: 0.9294 - loss: 0.2559 - val\_accuracy: 0.9458 - val\_loss: 0.2067  
 Epoch 42/100  
 1688/1688                    5s 2ms/step -  
 accuracy: 0.9276 - loss: 0.2579 - val\_accuracy: 0.9467 - val\_loss: 0.2047  
 Epoch 43/100  
 1688/1688                    5s 2ms/step -  
 accuracy: 0.9296 - loss: 0.2554 - val\_accuracy: 0.9465 - val\_loss: 0.2028  
 Epoch 44/100  
 1688/1688                    4s 2ms/step -  
 accuracy: 0.9307 - loss: 0.2496 - val\_accuracy: 0.9460 - val\_loss: 0.2011  
 Epoch 45/100  
 1688/1688                    5s 2ms/step -  
 accuracy: 0.9326 - loss: 0.2463 - val\_accuracy: 0.9468 - val\_loss: 0.1993  
 Epoch 46/100  
 1688/1688                    5s 2ms/step -  
 accuracy: 0.9316 - loss: 0.2479 - val\_accuracy: 0.9477 - val\_loss: 0.1975  
 Epoch 47/100  
 1688/1688                    5s 2ms/step -  
 accuracy: 0.9320 - loss: 0.2451 - val\_accuracy: 0.9480 - val\_loss: 0.1961  
 Epoch 48/100  
 1688/1688                    5s 2ms/step -

accuracy: 0.9341 - loss: 0.2389 - val\_accuracy: 0.9480 - val\_loss: 0.1943  
 Epoch 49/100  
 1688/1688 4s 2ms/step -  
 accuracy: 0.9345 - loss: 0.2380 - val\_accuracy: 0.9485 - val\_loss: 0.1928  
 Epoch 50/100  
 1688/1688 5s 2ms/step -  
 accuracy: 0.9341 - loss: 0.2370 - val\_accuracy: 0.9492 - val\_loss: 0.1914  
 Epoch 51/100  
 1688/1688 3s 2ms/step -  
 accuracy: 0.9335 - loss: 0.2379 - val\_accuracy: 0.9490 - val\_loss: 0.1902  
 Epoch 52/100  
 1688/1688 4s 2ms/step -  
 accuracy: 0.9351 - loss: 0.2340 - val\_accuracy: 0.9495 - val\_loss: 0.1889  
 Epoch 53/100  
 1688/1688 5s 2ms/step -  
 accuracy: 0.9361 - loss: 0.2307 - val\_accuracy: 0.9503 - val\_loss: 0.1872  
 Epoch 54/100  
 1688/1688 3s 2ms/step -  
 accuracy: 0.9358 - loss: 0.2297 - val\_accuracy: 0.9508 - val\_loss: 0.1860  
 Epoch 55/100  
 1688/1688 4s 2ms/step -  
 accuracy: 0.9364 - loss: 0.2305 - val\_accuracy: 0.9510 - val\_loss: 0.1846  
 Epoch 56/100  
 1688/1688 3s 2ms/step -  
 accuracy: 0.9380 - loss: 0.2257 - val\_accuracy: 0.9517 - val\_loss: 0.1832  
 Epoch 57/100  
 1688/1688 5s 2ms/step -  
 accuracy: 0.9372 - loss: 0.2281 - val\_accuracy: 0.9517 - val\_loss: 0.1820  
 Epoch 58/100  
 1688/1688 5s 2ms/step -  
 accuracy: 0.9390 - loss: 0.2247 - val\_accuracy: 0.9517 - val\_loss: 0.1810  
 Epoch 59/100  
 1688/1688 3s 2ms/step -  
 accuracy: 0.9401 - loss: 0.2185 - val\_accuracy: 0.9523 - val\_loss: 0.1797  
 Epoch 60/100  
 1688/1688 6s 2ms/step -  
 accuracy: 0.9383 - loss: 0.2217 - val\_accuracy: 0.9525 - val\_loss: 0.1784  
 Epoch 61/100  
 1688/1688 3s 2ms/step -  
 accuracy: 0.9393 - loss: 0.2176 - val\_accuracy: 0.9518 - val\_loss: 0.1776  
 Epoch 62/100  
 1688/1688 5s 2ms/step -  
 accuracy: 0.9405 - loss: 0.2180 - val\_accuracy: 0.9528 - val\_loss: 0.1764  
 Epoch 63/100  
 1688/1688 4s 2ms/step -  
 accuracy: 0.9428 - loss: 0.2080 - val\_accuracy: 0.9520 - val\_loss: 0.1755  
 Epoch 64/100  
 1688/1688 3s 2ms/step -

accuracy: 0.9436 - loss: 0.2083 - val\_accuracy: 0.9528 - val\_loss: 0.1744  
 Epoch 65/100  
 1688/1688                    5s 2ms/step -  
 accuracy: 0.9406 - loss: 0.2127 - val\_accuracy: 0.9532 - val\_loss: 0.1732  
 Epoch 66/100  
 1688/1688                    5s 2ms/step -  
 accuracy: 0.9422 - loss: 0.2104 - val\_accuracy: 0.9525 - val\_loss: 0.1724  
 Epoch 67/100  
 1688/1688                    5s 2ms/step -  
 accuracy: 0.9419 - loss: 0.2092 - val\_accuracy: 0.9532 - val\_loss: 0.1714  
 Epoch 68/100  
 1688/1688                    6s 2ms/step -  
 accuracy: 0.9424 - loss: 0.2076 - val\_accuracy: 0.9543 - val\_loss: 0.1704  
 Epoch 69/100  
 1688/1688                    5s 2ms/step -  
 accuracy: 0.9442 - loss: 0.2053 - val\_accuracy: 0.9552 - val\_loss: 0.1696  
 Epoch 70/100  
 1688/1688                    6s 2ms/step -  
 accuracy: 0.9432 - loss: 0.2050 - val\_accuracy: 0.9540 - val\_loss: 0.1688  
 Epoch 71/100  
 1688/1688                    4s 2ms/step -  
 accuracy: 0.9454 - loss: 0.2000 - val\_accuracy: 0.9550 - val\_loss: 0.1678  
 Epoch 72/100  
 1688/1688                    6s 2ms/step -  
 accuracy: 0.9439 - loss: 0.2022 - val\_accuracy: 0.9538 - val\_loss: 0.1674  
 Epoch 73/100  
 1688/1688                    4s 2ms/step -  
 accuracy: 0.9431 - loss: 0.2007 - val\_accuracy: 0.9542 - val\_loss: 0.1663  
 Epoch 74/100  
 1688/1688                    3s 2ms/step -  
 accuracy: 0.9427 - loss: 0.2042 - val\_accuracy: 0.9558 - val\_loss: 0.1654  
 Epoch 75/100  
 1688/1688                    3s 2ms/step -  
 accuracy: 0.9437 - loss: 0.2016 - val\_accuracy: 0.9562 - val\_loss: 0.1643  
 Epoch 76/100  
 1688/1688                    5s 2ms/step -  
 accuracy: 0.9432 - loss: 0.2002 - val\_accuracy: 0.9558 - val\_loss: 0.1639  
 Epoch 77/100  
 1688/1688                    5s 2ms/step -  
 accuracy: 0.9453 - loss: 0.1972 - val\_accuracy: 0.9567 - val\_loss: 0.1630  
 Epoch 78/100  
 1688/1688                    6s 2ms/step -  
 accuracy: 0.9479 - loss: 0.1910 - val\_accuracy: 0.9567 - val\_loss: 0.1624  
 Epoch 79/100  
 1688/1688                    4s 2ms/step -  
 accuracy: 0.9467 - loss: 0.1942 - val\_accuracy: 0.9570 - val\_loss: 0.1617  
 Epoch 80/100  
 1688/1688                    6s 2ms/step -

accuracy: 0.9450 - loss: 0.1957 - val\_accuracy: 0.9575 - val\_loss: 0.1606  
 Epoch 81/100  
 1688/1688 4s 2ms/step -  
 accuracy: 0.9460 - loss: 0.1949 - val\_accuracy: 0.9573 - val\_loss: 0.1599  
 Epoch 82/100  
 1688/1688 3s 2ms/step -  
 accuracy: 0.9435 - loss: 0.1985 - val\_accuracy: 0.9572 - val\_loss: 0.1595  
 Epoch 83/100  
 1688/1688 5s 2ms/step -  
 accuracy: 0.9451 - loss: 0.1967 - val\_accuracy: 0.9577 - val\_loss: 0.1586  
 Epoch 84/100  
 1688/1688 5s 2ms/step -  
 accuracy: 0.9471 - loss: 0.1901 - val\_accuracy: 0.9587 - val\_loss: 0.1579  
 Epoch 85/100  
 1688/1688 5s 2ms/step -  
 accuracy: 0.9472 - loss: 0.1886 - val\_accuracy: 0.9573 - val\_loss: 0.1575  
 Epoch 86/100  
 1688/1688 5s 2ms/step -  
 accuracy: 0.9478 - loss: 0.1849 - val\_accuracy: 0.9582 - val\_loss: 0.1569  
 Epoch 87/100  
 1688/1688 3s 2ms/step -  
 accuracy: 0.9478 - loss: 0.1839 - val\_accuracy: 0.9583 - val\_loss: 0.1563  
 Epoch 88/100  
 1688/1688 5s 2ms/step -  
 accuracy: 0.9486 - loss: 0.1839 - val\_accuracy: 0.9585 - val\_loss: 0.1554  
 Epoch 89/100  
 1688/1688 5s 2ms/step -  
 accuracy: 0.9469 - loss: 0.1876 - val\_accuracy: 0.9587 - val\_loss: 0.1550  
 Epoch 90/100  
 1688/1688 3s 2ms/step -  
 accuracy: 0.9481 - loss: 0.1857 - val\_accuracy: 0.9582 - val\_loss: 0.1545  
 Epoch 91/100  
 1688/1688 5s 2ms/step -  
 accuracy: 0.9487 - loss: 0.1840 - val\_accuracy: 0.9582 - val\_loss: 0.1539  
 Epoch 92/100  
 1688/1688 3s 2ms/step -  
 accuracy: 0.9499 - loss: 0.1817 - val\_accuracy: 0.9578 - val\_loss: 0.1536  
 Epoch 93/100  
 1688/1688 4s 2ms/step -  
 accuracy: 0.9501 - loss: 0.1806 - val\_accuracy: 0.9585 - val\_loss: 0.1528  
 Epoch 94/100  
 1688/1688 4s 2ms/step -  
 accuracy: 0.9497 - loss: 0.1811 - val\_accuracy: 0.9587 - val\_loss: 0.1521  
 Epoch 95/100  
 1688/1688 3s 2ms/step -  
 accuracy: 0.9499 - loss: 0.1778 - val\_accuracy: 0.9597 - val\_loss: 0.1516  
 Epoch 96/100  
 1688/1688 4s 2ms/step -

accuracy: 0.9491 - loss: 0.1816 - val\_accuracy: 0.9588 - val\_loss: 0.1513  
 Epoch 97/100  
 1688/1688 3s 2ms/step -  
 accuracy: 0.9496 - loss: 0.1804 - val\_accuracy: 0.9590 - val\_loss: 0.1505  
 Epoch 98/100  
 1688/1688 3s 2ms/step -  
 accuracy: 0.9503 - loss: 0.1776 - val\_accuracy: 0.9595 - val\_loss: 0.1499  
 Epoch 99/100  
 1688/1688 5s 2ms/step -  
 accuracy: 0.9503 - loss: 0.1788 - val\_accuracy: 0.9588 - val\_loss: 0.1499  
 Epoch 100/100  
 1688/1688 3s 2ms/step -  
 accuracy: 0.9506 - loss: 0.1713 - val\_accuracy: 0.9593 - val\_loss: 0.1488  
 Restoring model weights from the end of the best epoch: 100.  
 Training with batch size 64 and learning rate 0.01  
 Epoch 1/100  
 844/844 4s 3ms/step -  
 accuracy: 0.8417 - loss: 0.5082 - val\_accuracy: 0.9510 - val\_loss: 0.1620  
 Epoch 2/100  
 844/844 3s 3ms/step -  
 accuracy: 0.9440 - loss: 0.1895 - val\_accuracy: 0.9568 - val\_loss: 0.1460  
 Epoch 3/100  
 844/844 2s 3ms/step -  
 accuracy: 0.9512 - loss: 0.1617 - val\_accuracy: 0.9627 - val\_loss: 0.1439  
 Epoch 4/100  
 844/844 2s 3ms/step -  
 accuracy: 0.9586 - loss: 0.1401 - val\_accuracy: 0.9638 - val\_loss: 0.1254  
 Epoch 5/100  
 844/844 3s 3ms/step -  
 accuracy: 0.9624 - loss: 0.1276 - val\_accuracy: 0.9575 - val\_loss: 0.1410  
 Epoch 6/100  
 844/844 3s 3ms/step -  
 accuracy: 0.9641 - loss: 0.1189 - val\_accuracy: 0.9565 - val\_loss: 0.1493  
 Epoch 7/100  
 844/844 2s 3ms/step -  
 accuracy: 0.9643 - loss: 0.1154 - val\_accuracy: 0.9602 - val\_loss: 0.1563  
 Epoch 7: early stopping  
 Restoring model weights from the end of the best epoch: 4.  
 Training with batch size 64 and learning rate 0.001  
 Epoch 1/100  
 844/844 3s 3ms/step -  
 accuracy: 0.7561 - loss: 0.8474 - val\_accuracy: 0.9447 - val\_loss: 0.2059  
 Epoch 2/100  
 844/844 2s 2ms/step -  
 accuracy: 0.9298 - loss: 0.2476 - val\_accuracy: 0.9577 - val\_loss: 0.1564  
 Epoch 3/100  
 844/844 2s 2ms/step -  
 accuracy: 0.9471 - loss: 0.1836 - val\_accuracy: 0.9625 - val\_loss: 0.1329



Epoch 4/100  
844/844 3s 3ms/step -  
accuracy: 0.9565 - loss: 0.1515 - val\_accuracy: 0.9628 - val\_loss: 0.1263  
Epoch 5/100  
844/844 2s 2ms/step -  
accuracy: 0.9627 - loss: 0.1292 - val\_accuracy: 0.9620 - val\_loss: 0.1265  
Epoch 6/100  
844/844 2s 2ms/step -  
accuracy: 0.9664 - loss: 0.1130 - val\_accuracy: 0.9705 - val\_loss: 0.1039  
Epoch 7/100  
844/844 3s 2ms/step -  
accuracy: 0.9717 - loss: 0.0971 - val\_accuracy: 0.9675 - val\_loss: 0.1064  
Epoch 8/100  
844/844 3s 3ms/step -  
accuracy: 0.9729 - loss: 0.0878 - val\_accuracy: 0.9677 - val\_loss: 0.1019  
Epoch 9/100  
844/844 3s 3ms/step -  
accuracy: 0.9776 - loss: 0.0773 - val\_accuracy: 0.9677 - val\_loss: 0.1034  
Epoch 10/100  
844/844 2s 3ms/step -  
accuracy: 0.9781 - loss: 0.0723 - val\_accuracy: 0.9673 - val\_loss: 0.1143  
Epoch 11/100  
844/844 3s 3ms/step -  
accuracy: 0.9791 - loss: 0.0707 - val\_accuracy: 0.9713 - val\_loss: 0.0970  
Epoch 12/100  
844/844 2s 2ms/step -  
accuracy: 0.9804 - loss: 0.0666 - val\_accuracy: 0.9700 - val\_loss: 0.1031  
Epoch 13/100  
844/844 3s 3ms/step -  
accuracy: 0.9811 - loss: 0.0618 - val\_accuracy: 0.9725 - val\_loss: 0.0946  
Epoch 14/100  
844/844 5s 2ms/step -  
accuracy: 0.9825 - loss: 0.0573 - val\_accuracy: 0.9720 - val\_loss: 0.0960  
Epoch 15/100  
844/844 3s 2ms/step -  
accuracy: 0.9842 - loss: 0.0520 - val\_accuracy: 0.9702 - val\_loss: 0.1044  
Epoch 16/100  
844/844 3s 3ms/step -  
accuracy: 0.9853 - loss: 0.0492 - val\_accuracy: 0.9715 - val\_loss: 0.1009  
Epoch 16: early stopping  
Restoring model weights from the end of the best epoch: 13.  
Training with batch size 64 and learning rate 0.0001  
Epoch 1/100  
844/844 3s 3ms/step -  
accuracy: 0.4248 - loss: 1.7688 - val\_accuracy: 0.8727 - val\_loss: 0.5855  
Epoch 2/100  
844/844 2s 3ms/step -  
accuracy: 0.8593 - loss: 0.5714 - val\_accuracy: 0.9130 - val\_loss: 0.3516

Epoch 3/100  
844/844 2s 3ms/step -  
accuracy: 0.8920 - loss: 0.4045 - val\_accuracy: 0.9235 - val\_loss: 0.2854

Epoch 4/100  
844/844 3s 3ms/step -  
accuracy: 0.9082 - loss: 0.3388 - val\_accuracy: 0.9333 - val\_loss: 0.2519

Epoch 5/100  
844/844 2s 2ms/step -  
accuracy: 0.9173 - loss: 0.3019 - val\_accuracy: 0.9368 - val\_loss: 0.2307

Epoch 6/100  
844/844 2s 2ms/step -  
accuracy: 0.9219 - loss: 0.2817 - val\_accuracy: 0.9420 - val\_loss: 0.2164

Epoch 7/100  
844/844 3s 2ms/step -  
accuracy: 0.9260 - loss: 0.2596 - val\_accuracy: 0.9457 - val\_loss: 0.2045

Epoch 8/100  
844/844 2s 2ms/step -  
accuracy: 0.9303 - loss: 0.2498 - val\_accuracy: 0.9465 - val\_loss: 0.1962

Epoch 9/100  
844/844 2s 3ms/step -  
accuracy: 0.9328 - loss: 0.2403 - val\_accuracy: 0.9502 - val\_loss: 0.1877

Epoch 10/100  
844/844 2s 3ms/step -  
accuracy: 0.9352 - loss: 0.2267 - val\_accuracy: 0.9527 - val\_loss: 0.1807

Epoch 11/100  
844/844 2s 2ms/step -  
accuracy: 0.9373 - loss: 0.2231 - val\_accuracy: 0.9543 - val\_loss: 0.1747

Epoch 12/100  
844/844 2s 2ms/step -  
accuracy: 0.9401 - loss: 0.2112 - val\_accuracy: 0.9555 - val\_loss: 0.1702

Epoch 13/100  
844/844 3s 2ms/step -  
accuracy: 0.9415 - loss: 0.2055 - val\_accuracy: 0.9567 - val\_loss: 0.1650

Epoch 14/100  
844/844 3s 3ms/step -  
accuracy: 0.9439 - loss: 0.2007 - val\_accuracy: 0.9583 - val\_loss: 0.1607

Epoch 15/100  
844/844 2s 2ms/step -  
accuracy: 0.9469 - loss: 0.1898 - val\_accuracy: 0.9592 - val\_loss: 0.1566

Epoch 16/100  
844/844 2s 2ms/step -  
accuracy: 0.9477 - loss: 0.1851 - val\_accuracy: 0.9597 - val\_loss: 0.1543

Epoch 17/100  
844/844 2s 3ms/step -  
accuracy: 0.9504 - loss: 0.1752 - val\_accuracy: 0.9598 - val\_loss: 0.1501

Epoch 18/100  
844/844 3s 3ms/step -  
accuracy: 0.9485 - loss: 0.1815 - val\_accuracy: 0.9608 - val\_loss: 0.1472

Epoch 19/100  
844/844                2s 3ms/step -  
accuracy: 0.9516 - loss: 0.1711 - val\_accuracy: 0.9603 - val\_loss: 0.1471  
Epoch 20/100  
844/844                2s 3ms/step -  
accuracy: 0.9524 - loss: 0.1654 - val\_accuracy: 0.9615 - val\_loss: 0.1427  
Epoch 21/100  
844/844                2s 2ms/step -  
accuracy: 0.9545 - loss: 0.1608 - val\_accuracy: 0.9607 - val\_loss: 0.1424  
Epoch 22/100  
844/844                2s 2ms/step -  
accuracy: 0.9542 - loss: 0.1595 - val\_accuracy: 0.9610 - val\_loss: 0.1388  
Epoch 23/100  
844/844                3s 2ms/step -  
accuracy: 0.9567 - loss: 0.1532 - val\_accuracy: 0.9638 - val\_loss: 0.1372  
Epoch 24/100  
844/844                2s 3ms/step -  
accuracy: 0.9558 - loss: 0.1528 - val\_accuracy: 0.9623 - val\_loss: 0.1368  
Epoch 25/100  
844/844                2s 3ms/step -  
accuracy: 0.9553 - loss: 0.1519 - val\_accuracy: 0.9632 - val\_loss: 0.1341  
Epoch 26/100  
844/844                2s 2ms/step -  
accuracy: 0.9598 - loss: 0.1404 - val\_accuracy: 0.9628 - val\_loss: 0.1341  
Epoch 27/100  
844/844                2s 2ms/step -  
accuracy: 0.9609 - loss: 0.1382 - val\_accuracy: 0.9650 - val\_loss: 0.1314  
Epoch 28/100  
844/844                3s 2ms/step -  
accuracy: 0.9595 - loss: 0.1424 - val\_accuracy: 0.9650 - val\_loss: 0.1303  
Epoch 29/100  
844/844                3s 3ms/step -  
accuracy: 0.9621 - loss: 0.1340 - val\_accuracy: 0.9655 - val\_loss: 0.1287  
Epoch 30/100  
844/844                5s 2ms/step -  
accuracy: 0.9622 - loss: 0.1350 - val\_accuracy: 0.9655 - val\_loss: 0.1276  
Epoch 31/100  
844/844                2s 2ms/step -  
accuracy: 0.9612 - loss: 0.1359 - val\_accuracy: 0.9653 - val\_loss: 0.1262  
Epoch 32/100  
844/844                2s 2ms/step -  
accuracy: 0.9618 - loss: 0.1323 - val\_accuracy: 0.9657 - val\_loss: 0.1252  
Epoch 33/100  
844/844                2s 3ms/step -  
accuracy: 0.9653 - loss: 0.1230 - val\_accuracy: 0.9652 - val\_loss: 0.1242  
Epoch 34/100  
844/844                2s 2ms/step -  
accuracy: 0.9638 - loss: 0.1254 - val\_accuracy: 0.9665 - val\_loss: 0.1229

Epoch 35/100  
844/844 3s 2ms/step -  
accuracy: 0.9654 - loss: 0.1222 - val\_accuracy: 0.9652 - val\_loss: 0.1252  
Epoch 36/100  
844/844 2s 2ms/step -  
accuracy: 0.9653 - loss: 0.1234 - val\_accuracy: 0.9663 - val\_loss: 0.1221  
Epoch 37/100  
844/844 3s 2ms/step -  
accuracy: 0.9656 - loss: 0.1224 - val\_accuracy: 0.9673 - val\_loss: 0.1218  
Epoch 38/100  
844/844 3s 3ms/step -  
accuracy: 0.9675 - loss: 0.1148 - val\_accuracy: 0.9678 - val\_loss: 0.1201  
Epoch 39/100  
844/844 2s 2ms/step -  
accuracy: 0.9670 - loss: 0.1132 - val\_accuracy: 0.9670 - val\_loss: 0.1207  
Epoch 40/100  
844/844 2s 2ms/step -  
accuracy: 0.9674 - loss: 0.1142 - val\_accuracy: 0.9675 - val\_loss: 0.1188  
Epoch 41/100  
844/844 2s 2ms/step -  
accuracy: 0.9670 - loss: 0.1142 - val\_accuracy: 0.9668 - val\_loss: 0.1199  
Epoch 42/100  
844/844 3s 3ms/step -  
accuracy: 0.9676 - loss: 0.1118 - val\_accuracy: 0.9677 - val\_loss: 0.1198  
Epoch 43/100  
844/844 2s 3ms/step -  
accuracy: 0.9683 - loss: 0.1118 - val\_accuracy: 0.9685 - val\_loss: 0.1180  
Epoch 44/100  
844/844 2s 2ms/step -  
accuracy: 0.9685 - loss: 0.1070 - val\_accuracy: 0.9678 - val\_loss: 0.1188  
Epoch 45/100  
844/844 3s 2ms/step -  
accuracy: 0.9702 - loss: 0.1038 - val\_accuracy: 0.9663 - val\_loss: 0.1189  
Epoch 46/100  
844/844 2s 2ms/step -  
accuracy: 0.9695 - loss: 0.1042 - val\_accuracy: 0.9688 - val\_loss: 0.1170  
Epoch 47/100  
844/844 3s 3ms/step -  
accuracy: 0.9698 - loss: 0.1017 - val\_accuracy: 0.9683 - val\_loss: 0.1175  
Epoch 48/100  
844/844 2s 2ms/step -  
accuracy: 0.9708 - loss: 0.1030 - val\_accuracy: 0.9678 - val\_loss: 0.1166  
Epoch 49/100  
844/844 3s 2ms/step -  
accuracy: 0.9717 - loss: 0.0981 - val\_accuracy: 0.9688 - val\_loss: 0.1159  
Epoch 50/100  
844/844 2s 2ms/step -  
accuracy: 0.9720 - loss: 0.0967 - val\_accuracy: 0.9682 - val\_loss: 0.1156

Epoch 51/100  
844/844 3s 2ms/step -  
accuracy: 0.9722 - loss: 0.0967 - val\_accuracy: 0.9678 - val\_loss: 0.1157  
Epoch 52/100  
844/844 3s 3ms/step -  
accuracy: 0.9721 - loss: 0.0975 - val\_accuracy: 0.9695 - val\_loss: 0.1145  
Epoch 53/100  
844/844 2s 2ms/step -  
accuracy: 0.9710 - loss: 0.0963 - val\_accuracy: 0.9690 - val\_loss: 0.1145  
Epoch 54/100  
844/844 2s 2ms/step -  
accuracy: 0.9733 - loss: 0.0918 - val\_accuracy: 0.9692 - val\_loss: 0.1146  
Epoch 55/100  
844/844 3s 2ms/step -  
accuracy: 0.9717 - loss: 0.0969 - val\_accuracy: 0.9683 - val\_loss: 0.1157  
Epoch 55: early stopping  
Restoring model weights from the end of the best epoch: 52.  
Training with batch size 64 and learning rate 1e-05  
Epoch 1/100  
844/844 4s 3ms/step -  
accuracy: 0.1760 - loss: 2.2626 - val\_accuracy: 0.3185 - val\_loss: 2.0967  
Epoch 2/100  
844/844 2s 2ms/step -  
accuracy: 0.3469 - loss: 2.0500 - val\_accuracy: 0.4382 - val\_loss: 1.8614  
Epoch 3/100  
844/844 2s 2ms/step -  
accuracy: 0.4417 - loss: 1.8275 - val\_accuracy: 0.5177 - val\_loss: 1.6438  
Epoch 4/100  
844/844 3s 2ms/step -  
accuracy: 0.5188 - loss: 1.6206 - val\_accuracy: 0.6133 - val\_loss: 1.4348  
Epoch 5/100  
844/844 2s 2ms/step -  
accuracy: 0.6034 - loss: 1.4206 - val\_accuracy: 0.6767 - val\_loss: 1.2486  
Epoch 6/100  
844/844 3s 3ms/step -  
accuracy: 0.6552 - loss: 1.2633 - val\_accuracy: 0.7225 - val\_loss: 1.0978  
Epoch 7/100  
844/844 2s 2ms/step -  
accuracy: 0.7011 - loss: 1.1239 - val\_accuracy: 0.7602 - val\_loss: 0.9729  
Epoch 8/100  
844/844 3s 2ms/step -  
accuracy: 0.7372 - loss: 1.0060 - val\_accuracy: 0.7902 - val\_loss: 0.8705  
Epoch 9/100  
844/844 2s 2ms/step -  
accuracy: 0.7596 - loss: 0.9187 - val\_accuracy: 0.8158 - val\_loss: 0.7864  
Epoch 10/100  
844/844 3s 2ms/step -  
accuracy: 0.7859 - loss: 0.8275 - val\_accuracy: 0.8337 - val\_loss: 0.7177

Epoch 11/100  
844/844                3s 3ms/step -  
accuracy: 0.8011 - loss: 0.7706 - val\_accuracy: 0.8475 - val\_loss: 0.6605

Epoch 12/100  
844/844                2s 2ms/step -  
accuracy: 0.8141 - loss: 0.7188 - val\_accuracy: 0.8575 - val\_loss: 0.6126

Epoch 13/100  
844/844                2s 2ms/step -  
accuracy: 0.8281 - loss: 0.6717 - val\_accuracy: 0.8667 - val\_loss: 0.5726

Epoch 14/100  
844/844                2s 2ms/step -  
accuracy: 0.8383 - loss: 0.6334 - val\_accuracy: 0.8753 - val\_loss: 0.5383

Epoch 15/100  
844/844                3s 2ms/step -  
accuracy: 0.8454 - loss: 0.5963 - val\_accuracy: 0.8827 - val\_loss: 0.5089

Epoch 16/100  
844/844                3s 3ms/step -  
accuracy: 0.8534 - loss: 0.5696 - val\_accuracy: 0.8867 - val\_loss: 0.4837

Epoch 17/100  
844/844                2s 2ms/step -  
accuracy: 0.8578 - loss: 0.5459 - val\_accuracy: 0.8913 - val\_loss: 0.4613

Epoch 18/100  
844/844                2s 2ms/step -  
accuracy: 0.8632 - loss: 0.5262 - val\_accuracy: 0.8945 - val\_loss: 0.4417

Epoch 19/100  
844/844                2s 2ms/step -  
accuracy: 0.8700 - loss: 0.4981 - val\_accuracy: 0.8983 - val\_loss: 0.4245

Epoch 20/100  
844/844                2s 2ms/step -  
accuracy: 0.8741 - loss: 0.4826 - val\_accuracy: 0.9003 - val\_loss: 0.4095

Epoch 21/100  
844/844                3s 3ms/step -  
accuracy: 0.8782 - loss: 0.4683 - val\_accuracy: 0.9043 - val\_loss: 0.3957

Epoch 22/100  
844/844                2s 2ms/step -  
accuracy: 0.8757 - loss: 0.4574 - val\_accuracy: 0.9077 - val\_loss: 0.3831

Epoch 23/100  
844/844                2s 2ms/step -  
accuracy: 0.8801 - loss: 0.4502 - val\_accuracy: 0.9082 - val\_loss: 0.3718

Epoch 24/100  
844/844                2s 2ms/step -  
accuracy: 0.8861 - loss: 0.4287 - val\_accuracy: 0.9088 - val\_loss: 0.3616

Epoch 25/100  
844/844                3s 2ms/step -  
accuracy: 0.8864 - loss: 0.4215 - val\_accuracy: 0.9128 - val\_loss: 0.3521

Epoch 26/100  
844/844                2s 3ms/step -  
accuracy: 0.8906 - loss: 0.4078 - val\_accuracy: 0.9145 - val\_loss: 0.3436



Epoch 27/100  
844/844 2s 2ms/step -  
accuracy: 0.8922 - loss: 0.4025 - val\_accuracy: 0.9158 - val\_loss: 0.3357  
Epoch 28/100  
844/844 2s 2ms/step -  
accuracy: 0.8949 - loss: 0.3920 - val\_accuracy: 0.9183 - val\_loss: 0.3282  
Epoch 29/100  
844/844 2s 2ms/step -  
accuracy: 0.8980 - loss: 0.3778 - val\_accuracy: 0.9192 - val\_loss: 0.3216  
Epoch 30/100  
844/844 2s 2ms/step -  
accuracy: 0.8976 - loss: 0.3780 - val\_accuracy: 0.9210 - val\_loss: 0.3153  
Epoch 31/100  
844/844 3s 3ms/step -  
accuracy: 0.8994 - loss: 0.3685 - val\_accuracy: 0.9220 - val\_loss: 0.3095  
Epoch 32/100  
844/844 2s 2ms/step -  
accuracy: 0.9021 - loss: 0.3646 - val\_accuracy: 0.9232 - val\_loss: 0.3042  
Epoch 33/100  
844/844 2s 2ms/step -  
accuracy: 0.9042 - loss: 0.3568 - val\_accuracy: 0.9242 - val\_loss: 0.2993  
Epoch 34/100  
844/844 2s 2ms/step -  
accuracy: 0.9034 - loss: 0.3568 - val\_accuracy: 0.9240 - val\_loss: 0.2944  
Epoch 35/100  
844/844 2s 2ms/step -  
accuracy: 0.9034 - loss: 0.3512 - val\_accuracy: 0.9257 - val\_loss: 0.2900  
Epoch 36/100  
844/844 2s 2ms/step -  
accuracy: 0.9055 - loss: 0.3441 - val\_accuracy: 0.9257 - val\_loss: 0.2859  
Epoch 37/100  
844/844 3s 3ms/step -  
accuracy: 0.9065 - loss: 0.3420 - val\_accuracy: 0.9270 - val\_loss: 0.2820  
Epoch 38/100  
844/844 2s 2ms/step -  
accuracy: 0.9082 - loss: 0.3348 - val\_accuracy: 0.9285 - val\_loss: 0.2783  
Epoch 39/100  
844/844 2s 2ms/step -  
accuracy: 0.9077 - loss: 0.3304 - val\_accuracy: 0.9290 - val\_loss: 0.2750  
Epoch 40/100  
844/844 3s 2ms/step -  
accuracy: 0.9113 - loss: 0.3219 - val\_accuracy: 0.9288 - val\_loss: 0.2717  
Epoch 41/100  
844/844 3s 2ms/step -  
accuracy: 0.9109 - loss: 0.3205 - val\_accuracy: 0.9300 - val\_loss: 0.2686  
Epoch 42/100  
844/844 2s 2ms/step -  
accuracy: 0.9135 - loss: 0.3154 - val\_accuracy: 0.9318 - val\_loss: 0.2658

Epoch 43/100  
844/844 2s 2ms/step -  
accuracy: 0.9122 - loss: 0.3183 - val\_accuracy: 0.9315 - val\_loss: 0.2630  
Epoch 44/100  
844/844 3s 2ms/step -  
accuracy: 0.9150 - loss: 0.3087 - val\_accuracy: 0.9323 - val\_loss: 0.2603  
Epoch 45/100  
844/844 2s 2ms/step -  
accuracy: 0.9145 - loss: 0.3112 - val\_accuracy: 0.9333 - val\_loss: 0.2579  
Epoch 46/100  
844/844 3s 3ms/step -  
accuracy: 0.9155 - loss: 0.3011 - val\_accuracy: 0.9335 - val\_loss: 0.2555  
Epoch 47/100  
844/844 2s 2ms/step -  
accuracy: 0.9179 - loss: 0.2985 - val\_accuracy: 0.9333 - val\_loss: 0.2531  
Epoch 48/100  
844/844 2s 2ms/step -  
accuracy: 0.9154 - loss: 0.3022 - val\_accuracy: 0.9342 - val\_loss: 0.2510  
Epoch 49/100  
844/844 3s 2ms/step -  
accuracy: 0.9179 - loss: 0.2958 - val\_accuracy: 0.9352 - val\_loss: 0.2489  
Epoch 50/100  
844/844 2s 2ms/step -  
accuracy: 0.9164 - loss: 0.2982 - val\_accuracy: 0.9350 - val\_loss: 0.2468  
Epoch 51/100  
844/844 3s 3ms/step -  
accuracy: 0.9187 - loss: 0.2970 - val\_accuracy: 0.9355 - val\_loss: 0.2449  
Epoch 52/100  
844/844 2s 2ms/step -  
accuracy: 0.9201 - loss: 0.2880 - val\_accuracy: 0.9362 - val\_loss: 0.2432  
Epoch 53/100  
844/844 3s 2ms/step -  
accuracy: 0.9205 - loss: 0.2887 - val\_accuracy: 0.9367 - val\_loss: 0.2413  
Epoch 54/100  
844/844 2s 2ms/step -  
accuracy: 0.9213 - loss: 0.2824 - val\_accuracy: 0.9370 - val\_loss: 0.2397  
Epoch 55/100  
844/844 2s 3ms/step -  
accuracy: 0.9217 - loss: 0.2813 - val\_accuracy: 0.9370 - val\_loss: 0.2381  
Epoch 56/100  
844/844 3s 3ms/step -  
accuracy: 0.9234 - loss: 0.2796 - val\_accuracy: 0.9368 - val\_loss: 0.2364  
Epoch 57/100  
844/844 5s 2ms/step -  
accuracy: 0.9230 - loss: 0.2774 - val\_accuracy: 0.9370 - val\_loss: 0.2349  
Epoch 58/100  
844/844 3s 2ms/step -  
accuracy: 0.9231 - loss: 0.2777 - val\_accuracy: 0.9373 - val\_loss: 0.2336

Epoch 59/100  
844/844 2s 2ms/step -  
accuracy: 0.9225 - loss: 0.2785 - val\_accuracy: 0.9368 - val\_loss: 0.2321  
Epoch 60/100  
844/844 3s 3ms/step -  
accuracy: 0.9249 - loss: 0.2713 - val\_accuracy: 0.9380 - val\_loss: 0.2307  
Epoch 61/100  
844/844 2s 2ms/step -  
accuracy: 0.9237 - loss: 0.2744 - val\_accuracy: 0.9387 - val\_loss: 0.2294  
Epoch 62/100  
844/844 2s 2ms/step -  
accuracy: 0.9254 - loss: 0.2705 - val\_accuracy: 0.9380 - val\_loss: 0.2281  
Epoch 63/100  
844/844 3s 3ms/step -  
accuracy: 0.9239 - loss: 0.2697 - val\_accuracy: 0.9388 - val\_loss: 0.2269  
Epoch 64/100  
844/844 2s 2ms/step -  
accuracy: 0.9282 - loss: 0.2579 - val\_accuracy: 0.9392 - val\_loss: 0.2256  
Epoch 65/100  
844/844 3s 3ms/step -  
accuracy: 0.9244 - loss: 0.2706 - val\_accuracy: 0.9390 - val\_loss: 0.2246  
Epoch 66/100  
844/844 2s 2ms/step -  
accuracy: 0.9245 - loss: 0.2683 - val\_accuracy: 0.9398 - val\_loss: 0.2234  
Epoch 67/100  
844/844 2s 2ms/step -  
accuracy: 0.9263 - loss: 0.2628 - val\_accuracy: 0.9408 - val\_loss: 0.2222  
Epoch 68/100  
844/844 3s 2ms/step -  
accuracy: 0.9265 - loss: 0.2603 - val\_accuracy: 0.9407 - val\_loss: 0.2211  
Epoch 69/100  
844/844 3s 2ms/step -  
accuracy: 0.9263 - loss: 0.2627 - val\_accuracy: 0.9415 - val\_loss: 0.2200  
Epoch 70/100  
844/844 3s 3ms/step -  
accuracy: 0.9256 - loss: 0.2607 - val\_accuracy: 0.9418 - val\_loss: 0.2190  
Epoch 71/100  
844/844 2s 2ms/step -  
accuracy: 0.9297 - loss: 0.2546 - val\_accuracy: 0.9422 - val\_loss: 0.2179  
Epoch 72/100  
844/844 2s 2ms/step -  
accuracy: 0.9295 - loss: 0.2499 - val\_accuracy: 0.9425 - val\_loss: 0.2171  
Epoch 73/100  
844/844 3s 2ms/step -  
accuracy: 0.9292 - loss: 0.2561 - val\_accuracy: 0.9427 - val\_loss: 0.2162  
Epoch 74/100  
844/844 3s 3ms/step -  
accuracy: 0.9303 - loss: 0.2513 - val\_accuracy: 0.9423 - val\_loss: 0.2153

Epoch 75/100  
844/844 4s 2ms/step -  
accuracy: 0.9297 - loss: 0.2511 - val\_accuracy: 0.9423 - val\_loss: 0.2143  
Epoch 76/100  
844/844 2s 2ms/step -  
accuracy: 0.9287 - loss: 0.2528 - val\_accuracy: 0.9430 - val\_loss: 0.2133  
Epoch 77/100  
844/844 3s 2ms/step -  
accuracy: 0.9279 - loss: 0.2569 - val\_accuracy: 0.9432 - val\_loss: 0.2127  
Epoch 78/100  
844/844 3s 3ms/step -  
accuracy: 0.9302 - loss: 0.2494 - val\_accuracy: 0.9437 - val\_loss: 0.2117  
Epoch 79/100  
844/844 2s 2ms/step -  
accuracy: 0.9317 - loss: 0.2429 - val\_accuracy: 0.9433 - val\_loss: 0.2108  
Epoch 80/100  
844/844 2s 2ms/step -  
accuracy: 0.9312 - loss: 0.2422 - val\_accuracy: 0.9437 - val\_loss: 0.2100  
Epoch 81/100  
844/844 3s 3ms/step -  
accuracy: 0.9325 - loss: 0.2391 - val\_accuracy: 0.9437 - val\_loss: 0.2092  
Epoch 82/100  
844/844 2s 2ms/step -  
accuracy: 0.9322 - loss: 0.2423 - val\_accuracy: 0.9437 - val\_loss: 0.2083  
Epoch 83/100  
844/844 3s 3ms/step -  
accuracy: 0.9308 - loss: 0.2476 - val\_accuracy: 0.9437 - val\_loss: 0.2076  
Epoch 84/100  
844/844 2s 2ms/step -  
accuracy: 0.9335 - loss: 0.2390 - val\_accuracy: 0.9435 - val\_loss: 0.2068  
Epoch 85/100  
844/844 3s 2ms/step -  
accuracy: 0.9323 - loss: 0.2392 - val\_accuracy: 0.9442 - val\_loss: 0.2061  
Epoch 86/100  
844/844 2s 2ms/step -  
accuracy: 0.9324 - loss: 0.2428 - val\_accuracy: 0.9443 - val\_loss: 0.2053  
Epoch 87/100  
844/844 3s 2ms/step -  
accuracy: 0.9335 - loss: 0.2387 - val\_accuracy: 0.9450 - val\_loss: 0.2046  
Epoch 88/100  
844/844 3s 3ms/step -  
accuracy: 0.9329 - loss: 0.2386 - val\_accuracy: 0.9447 - val\_loss: 0.2039  
Epoch 89/100  
844/844 2s 2ms/step -  
accuracy: 0.9330 - loss: 0.2390 - val\_accuracy: 0.9452 - val\_loss: 0.2032  
Epoch 90/100  
844/844 2s 2ms/step -  
accuracy: 0.9326 - loss: 0.2396 - val\_accuracy: 0.9455 - val\_loss: 0.2025

```


Epoch 91/100
844/844          2s 2ms/step -
accuracy: 0.9329 - loss: 0.2361 - val_accuracy: 0.9450 - val_loss: 0.2019
Epoch 92/100
844/844          3s 2ms/step -
accuracy: 0.9347 - loss: 0.2338 - val_accuracy: 0.9450 - val_loss: 0.2014
Epoch 93/100
844/844          3s 3ms/step -
accuracy: 0.9344 - loss: 0.2351 - val_accuracy: 0.9458 - val_loss: 0.2006
Epoch 94/100
844/844          5s 2ms/step -
accuracy: 0.9373 - loss: 0.2262 - val_accuracy: 0.9458 - val_loss: 0.2000
Epoch 95/100
844/844          3s 3ms/step -
accuracy: 0.9346 - loss: 0.2332 - val_accuracy: 0.9453 - val_loss: 0.1993
Epoch 96/100
844/844          2s 2ms/step -
accuracy: 0.9360 - loss: 0.2283 - val_accuracy: 0.9455 - val_loss: 0.1987
Epoch 97/100
844/844          3s 3ms/step -
accuracy: 0.9360 - loss: 0.2305 - val_accuracy: 0.9455 - val_loss: 0.1981
Epoch 98/100
844/844          2s 2ms/step -
accuracy: 0.9342 - loss: 0.2351 - val_accuracy: 0.9460 - val_loss: 0.1975
Epoch 99/100
844/844          2s 2ms/step -
accuracy: 0.9366 - loss: 0.2296 - val_accuracy: 0.9462 - val_loss: 0.1969
Epoch 100/100
844/844          2s 2ms/step -
accuracy: 0.9355 - loss: 0.2314 - val_accuracy: 0.9467 - val_loss: 0.1964
Restoring model weights from the end of the best epoch: 100.

```


```

[6]: # ratios of batch size to learning rate and plot results
ratios = [bs/lr for bs, lr, _ in results]
accuracies = [acc for _, _, acc in results]

```



```
[7]: #plot
plt.figure(figsize=(10, 5))
plt.scatter(ratios, accuracies, color='blue')
plt.title('Test Accuracy vs. Ratio of Batch Size to Learning Rate')
plt.xlabel('Ratio of Batch Size to Learning Rate')
plt.ylabel('Test Accuracy')
plt.grid(True)
plt.xscale('log')
plt.show()
```



**Observations:** - The charts display how varying the ratio of batch size to learning rate affects test accuracy in training ANNs for the MNIST dataset. - Larger ratios generally correlate with higher test accuracies, indicating that either increasing batch size or decreasing the learning rate can enhance model performance. - There's a clear peak in performance at mid-range ratios, suggesting an optimal balance between batch size and learning rate, beyond which performance may plateau or even decrease, as seen in the highest ratio values.

[ ]:

